

Interprocess Communication

Chester Rebeiro
IIT Madras

Virtual Memory View

- During execution, each process can only view its virtual addresses,
- It cannot
 - View another processes virtual address space
 - Determine the physical address mapping

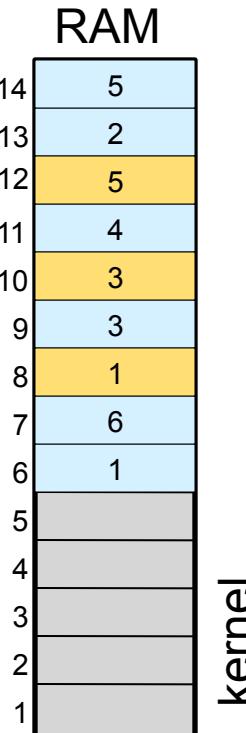
Executing
Process

Virtual Memory Map

6
5
4
3
2
1

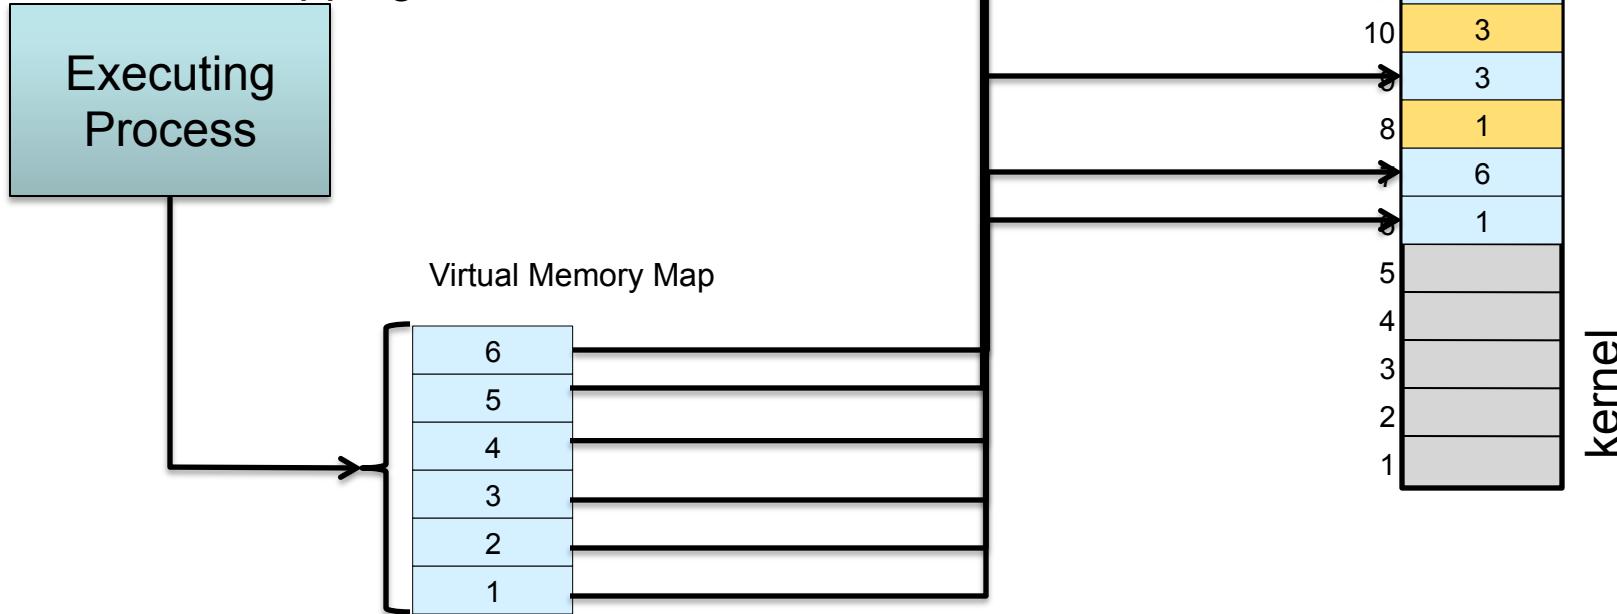
Virtual Memory Map

6
5
4
3
2
1



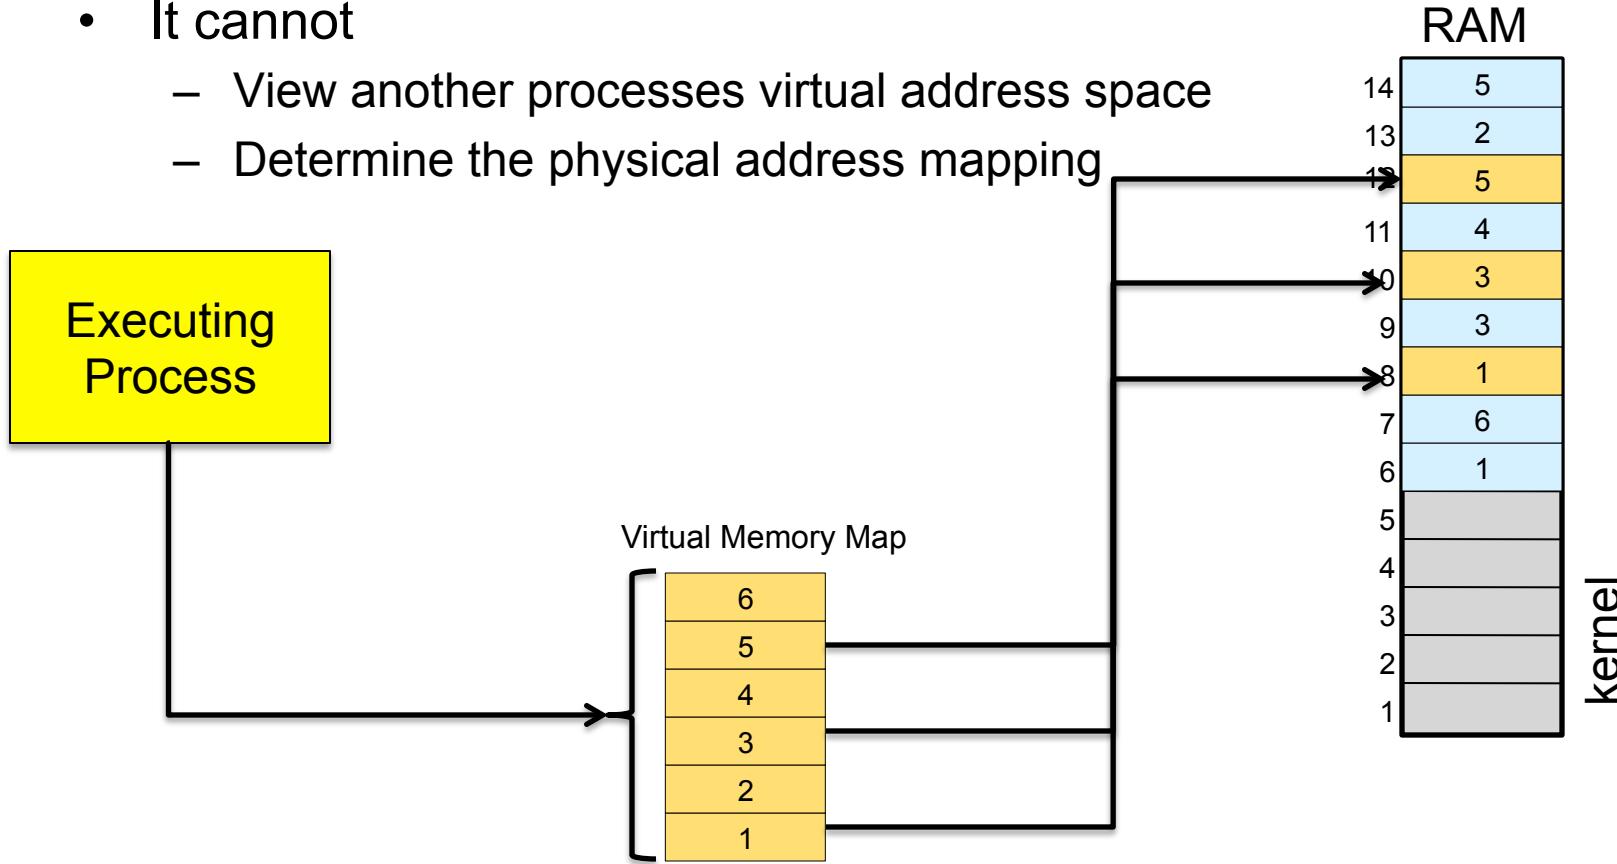
Virtual Memory View

- During execution, each process can only view its virtual addresses,
- It cannot
 - View another processes virtual address space
 - Determine the physical address mapping



Virtual Memory View

- During execution, each process can only view its virtual addresses,
- It cannot
 - View another processes virtual address space
 - Determine the physical address mapping

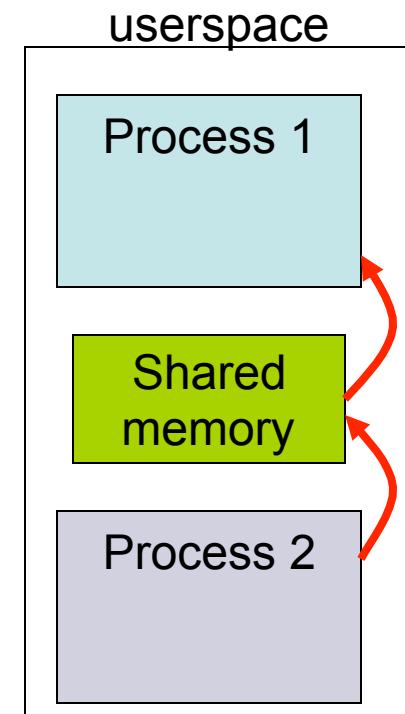


Inter Process Communication

- Advantages of Inter Process Communication (**IPC**)
 - Information sharing
 - Modularity/Convenience
- 3 ways
 - Shared memory
 - Message Passing
 - Signals

Shared Memory

- One process will create an area in RAM which the other process can access
- Both processes can access shared memory like a regular working memory
 - Reading/writing is like regular reading/writing
 - Fast
- **Limitation :** Error prone. Needs synchronization between processes



Shared Memory in Linux

- **int shmget (key, size, flags)**
 - Create a shared memory segment;
 - Returns ID of segment : **shmid**
 - **key** : unique identifier of the shared memory segment
 - **size** : size of the shared memory (rounded up to the **PAGE_SIZE**)
- **int shmat(shmid, addr, flags)**
 - **Attach** **shmid** shared memory to address space of the calling process
 - **addr** : pointer to the shared memory address space
- **int shmdt(shmid)**
 - **Detach** shared memory

Example

server.c

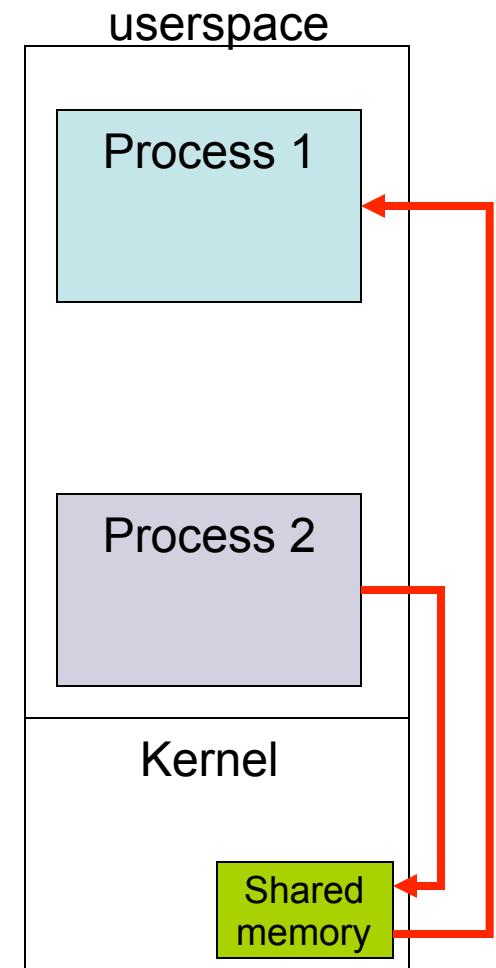
```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 #define SHMSIZE    27 /* Size of shared memory */
8
9 main()
10 {
11     char c;
12     int shmid;
13     key_t key;
14     char *shm, *s;
15
16     key = 5678; /* some key to uniquely identifies the shared memory */
17
18     /* Create the segment. */
19     if ((shmid = shmget(key, SHMSIZE, IPC_CREAT | 0666)) < 0) {
20         perror("shmget");
21         exit(1);
22     }
23
24     /* Attach the segment to our data space. */
25     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
26         perror("shmat");
27         exit(1);
28     }
29
30     /* Now put some things into the shared memory */
31     s = shm;
32     for (c = 'a'; c <= 'z'; c++)
33         *s++ = c;
34     *s = 0; /* end with a NULL termination */
35
36     /* Wait until the other process changes the first character
37      * to '*' the shared memory */
38     while (*shm != '*')
39         sleep(1);
40     exit(0);
41 }
```

client.c

```
1 #include <sys/types.h>
2 #include <sys/ipc.h>
3 #include <sys/shm.h>
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 #define SHMSIZE    27
8
9 main()
10 {
11     int shmid;
12     key_t key;
13     char *shm, *s;
14
15     /* We need to get the segment named "5678", created by the server
16     key = 5678;
17
18     /* Locate the segment. */
19     if ((shmid = shmget(key, SHMSIZE, 0666)) < 0) {
20         perror("shmget");
21         exit(1);
22     }
23
24     /* Attach the segment to our data space. */
25     if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
26         perror("shmat");
27         exit(1);
28     }
29
30     /* read what the server put in the memory. */
31     for (s = shm; *s != 0; s++)
32         putchar(*s);
33     putchar('\n');
34
35     /*
36      * Finally, change the first character of the
37      * segment to '*', indicating we have read
38      * the segment.
39      */
40     *shm = '*';
41
42     exit(0);
```

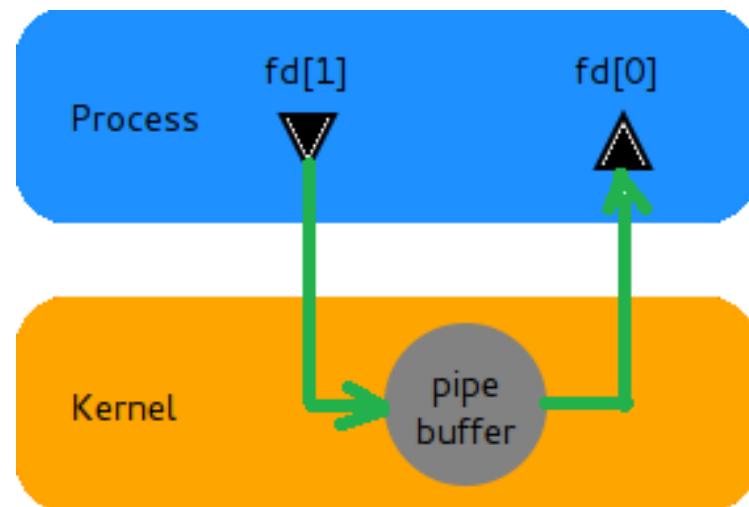
Message Passing

- Shared memory created in the kernel
- System calls such as **send** and **receive** used for communication
 - Cooperating : each send must have a receive
- **Advantage** : Explicit sharing, less error prone
- **Limitation** : Slow. Each call involves marshalling / demarshalling of information

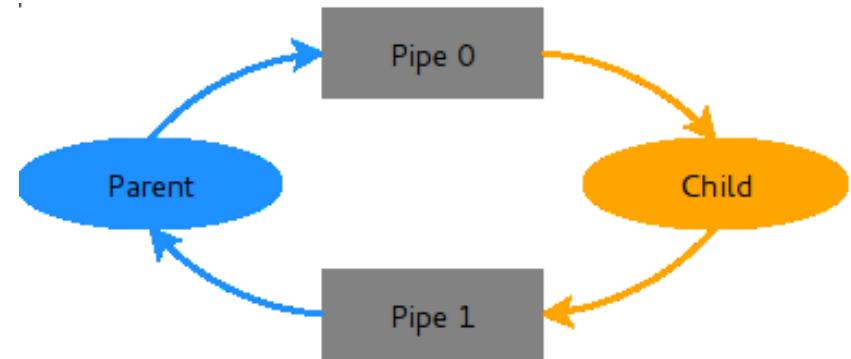
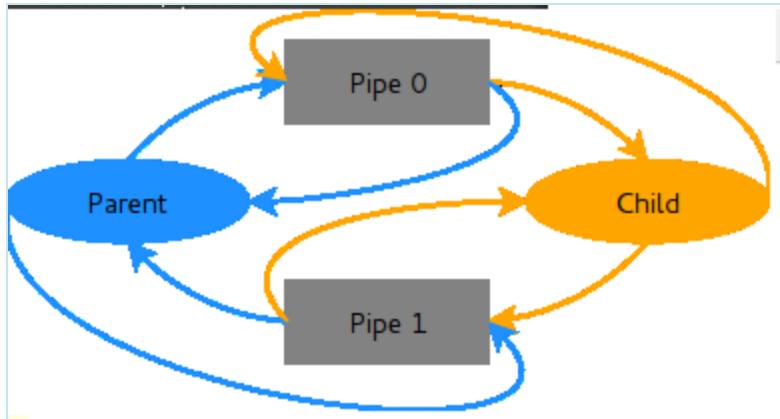


Pipes

- Always between parent and child
- Always unidirectional
- Accessed by two associated file descriptors:
 - fd[0] for reading from pipe
 - fd[1] for writing to the pipe



Pipes for two way communication



- Two pipes opened
pipe0 and pipe1
- Note the unnecessary
pipes
- Close the unnecessary
pipes

Example

(child process sending a string to parent)

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main(){
    int pipefd[2];
    int pid;
    char recv[32];

    pipe(pipefd);

    switch(pid=fork()) {
        case -1: perror("fork");
                    exit(1);
        case 0:          /* in child process */
                    close(pipefd[0]);
                    /* close unnecessary pipefd */
                    FILE *out = fdopen(pipefd[1], "w"); /* open pipe descriptor as stream */
                    fprintf(out, "Hello World\n");      /* write to out stream */
                    break;
        default:         /* in parent process */
                    close(pipefd[1]);
                    /* close unnecessary pipefd */
                    FILE *in = fdopen(pipefd[0], "r"); /* open descriptor as stream */
                    fscanf(in, "%s", recv);           /* read from in stream */
                    printf("%s", recv);
                    break;
    }
}
```

Signals

- Asynchronous unidirectional communication between processes
- Signals are a small integer
 - eg. 9: kill, 11: segmentation fault
- Send a signal to a process
 - `kill(pid, signum)`
- Process handler for a signal
 - `sighandler_t signal(signum, handler);`
 - Default if no handler defined

Synchronization

Chester Rebeiro
IIT Madras

Motivating Scenario

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

- Single core
 - Program 1 and program 2 are executing at the same time but sharing a single core



→CPU usage wrt time

Motivating Scenario

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

Shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

- What is the value of counter?
 - expected to be 5
 - but could also be 4 and 6

the general idea - when there are multiple accesses to shared memory by different processes/threads, there may be indeterminism in the results - the final result depends on the order in which the accesses to the memory are made

Motivating Scenario

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

Shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

R1 \leftarrow counter
R1 \leftarrow *R1* + 1
counter \leftarrow *R1*
R2 \leftarrow counter
R2 \leftarrow *R2* - 1
counter \leftarrow *R2*

R1 \leftarrow counter
R2 \leftarrow counter
R2 \leftarrow *R2* - 1
counter \leftarrow *R2*
R1 \leftarrow *R1* + 1
counter \leftarrow *R1*

R2 \leftarrow counter
R2 \leftarrow counter
R2 \leftarrow *R2* + 1
counter \leftarrow *R2*
R2 \leftarrow *R2* - 1
counter \leftarrow *R2*

context
switch

counter = 5

counter = 6

counter = 4

Race Conditions

- Race conditions
 - A situation where several processes access and manipulate the same data (*critical section*)
 - The outcome depends on the order in which the access take place
 - Prevent race conditions by synchronization
 - Ensure only one process at a time manipulates the critical data

```
{  
*  
*  
counter++  
*  
}
```

critical section

*No more than one
process should execute in
critical section at a time*

Race Conditions in Multicore

program 0

```
{  
    *  
    *  
    counter++  
    *  
}
```

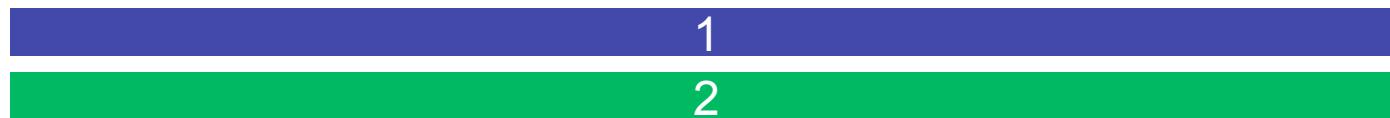
shared variable

```
int counter=5;
```

program 1

```
{  
    *  
    *  
    counter--  
    *  
}
```

- Multi core
 - Program 1 and program 2 are executing at the same time on different cores



→CPU usage wrt time

Critical Section

- Any solution should satisfy the following requirements

this is the entire point - more than one access - data race and indeterminism

- **Mutual Exclusion** : No more than one process in critical section at a given time
 - there must be no stalling of processes unnecessarily -
 - **Progress** : When no process is in the critical section, any process that requests entry into the critical section must be permitted without any delay
 - no process should have to wait forever to gain access to the critical section
 - **No starvation (bounded wait)**: There is an upper bound on the number of times a process enters the critical section, while another is waiting.

Locks and Unlocks

program 0

```
{  
    *  
    *  
    lock(L)  
    counter++  
    unlock(L)  
    *  
}
```

shared variable

```
int counter=5;  
lock_t L;
```

program 1

```
{  
    *  
    *  
    lock(L)  
    counter--  
    unlock(L)  
    *  
}
```

- **lock(L)** : acquire lock L exclusively
 - Only the process with L can access the critical section
- **unlock(L)** : release exclusive access to lock L
 - Permitting other processes to access the critical section

locks satisfy the three conditions for solutions to the race problem:

1. mutual exclusion - only a process holding the lock may gain access to the critical section
2. progress - when no process is holding the lock, the acquiring of the lock happens immediately and critical section is accessed.
3. bounded wait - there will be a limit imposed on how long one process can hold the lock when another is attempting to acquire.

When to have Locking?

- Single instructions by themselves are atomic
 - eg. add %eax, %ebx
- Multiple instructions need to be explicitly made atomic
 - Each piece of code in the OS must be checked if they need to be atomic

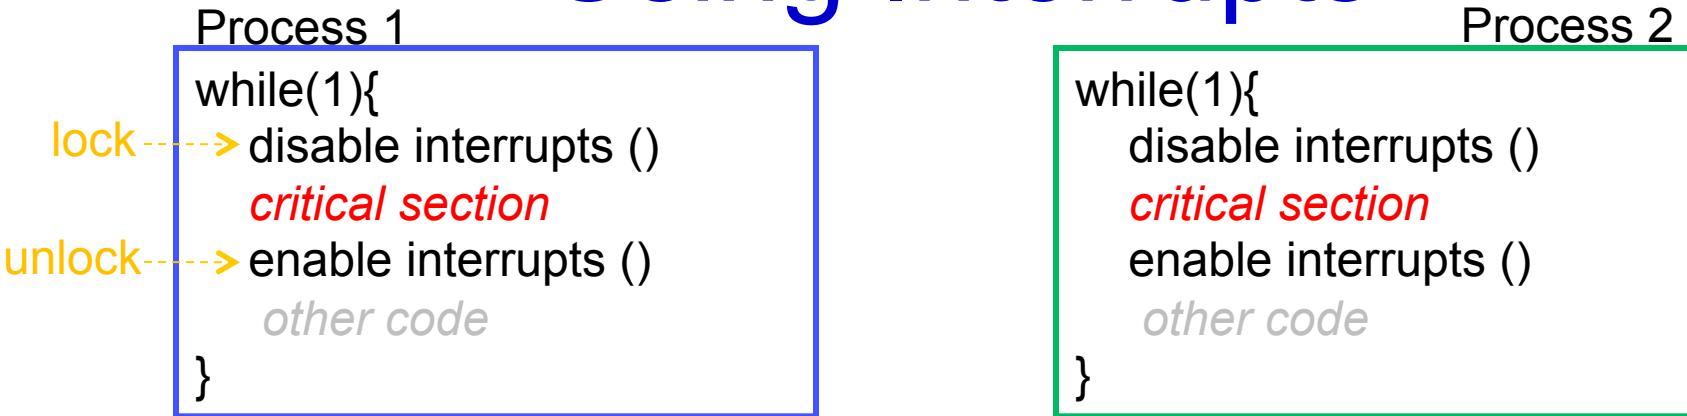
important pieces of code such as interrupt handlers, small critical sections etc need to be made atomic.

note that coming up with instructions that themselves are atomic is not optimal and would lead to complicated microarchitecture instead blocks of code may be covered in constructs that prevent other processes from executing that piece = all or nothing approach

How to Implement Locking (Software Solutions)

Chester Rebeiro
IIT Madras

Using Interrupts

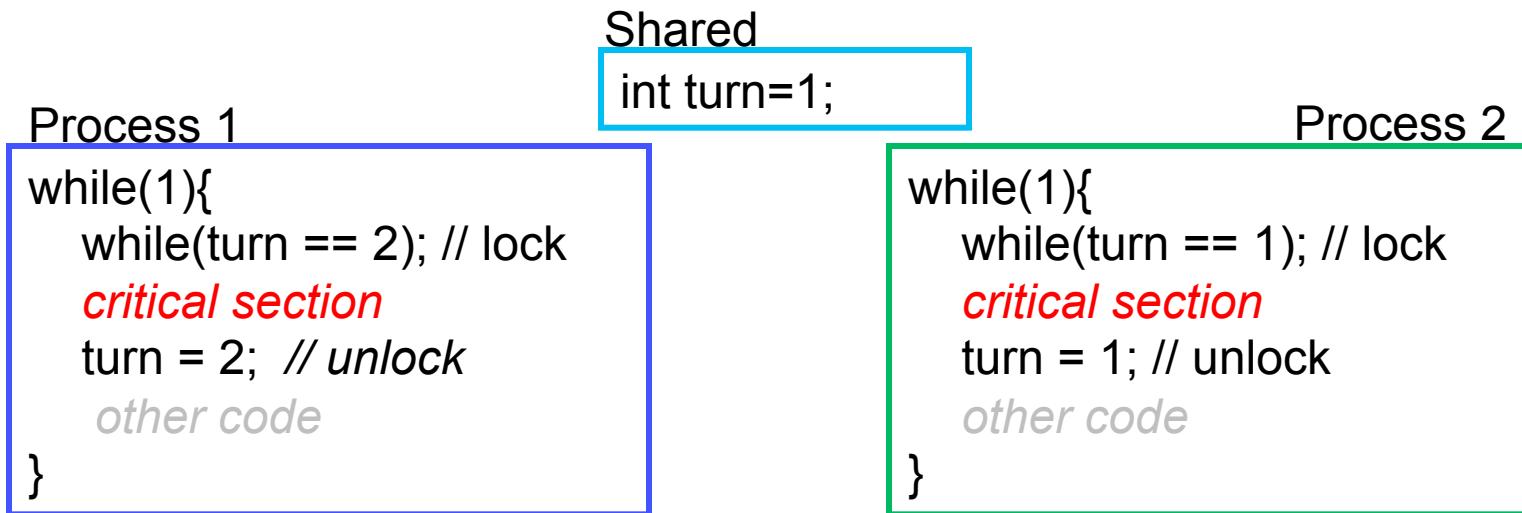


- **Simple**
 - When interrupts are disabled, context switches won't happen
- **Requires privileges**
 - User processes generally cannot disable interrupts
- **Not suited for multicore systems**

this can only be permitted for kernel code and on single core systems

this can also have other problems such as disabling interrupts for a long time can lead to missing important interrupt calls

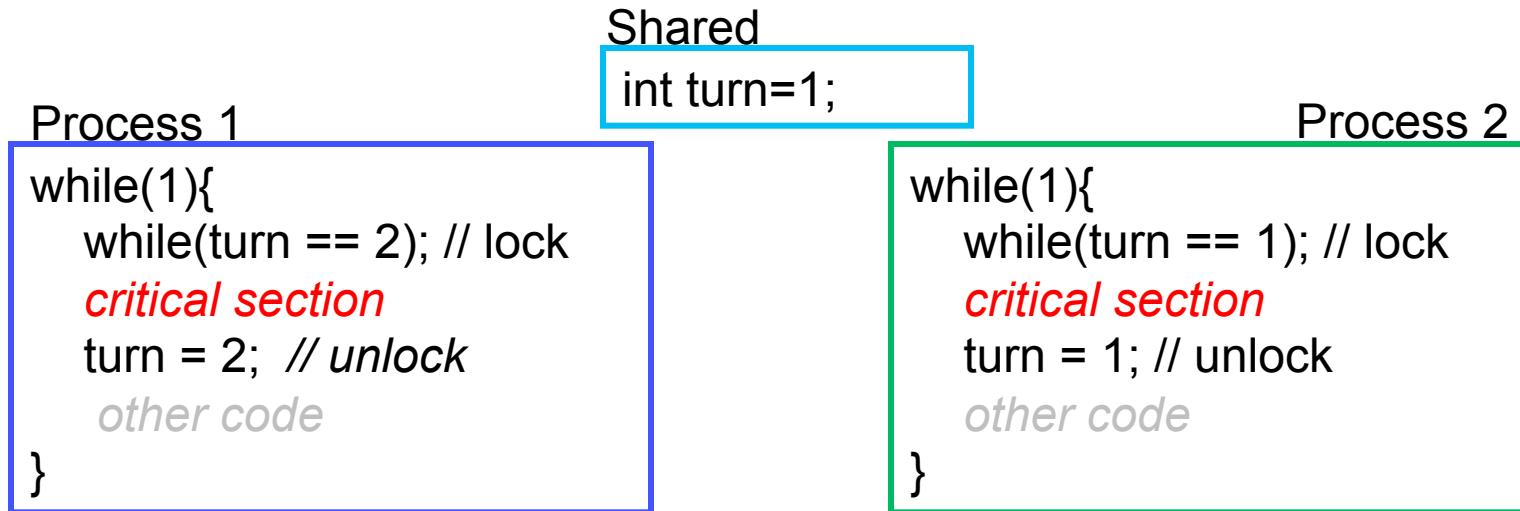
Software Solution (Attempt 1)



this is also weird in that it does not satisfy the PROGRESS requirement - if process 1 did not run, process 2 will have to wait forever to enter

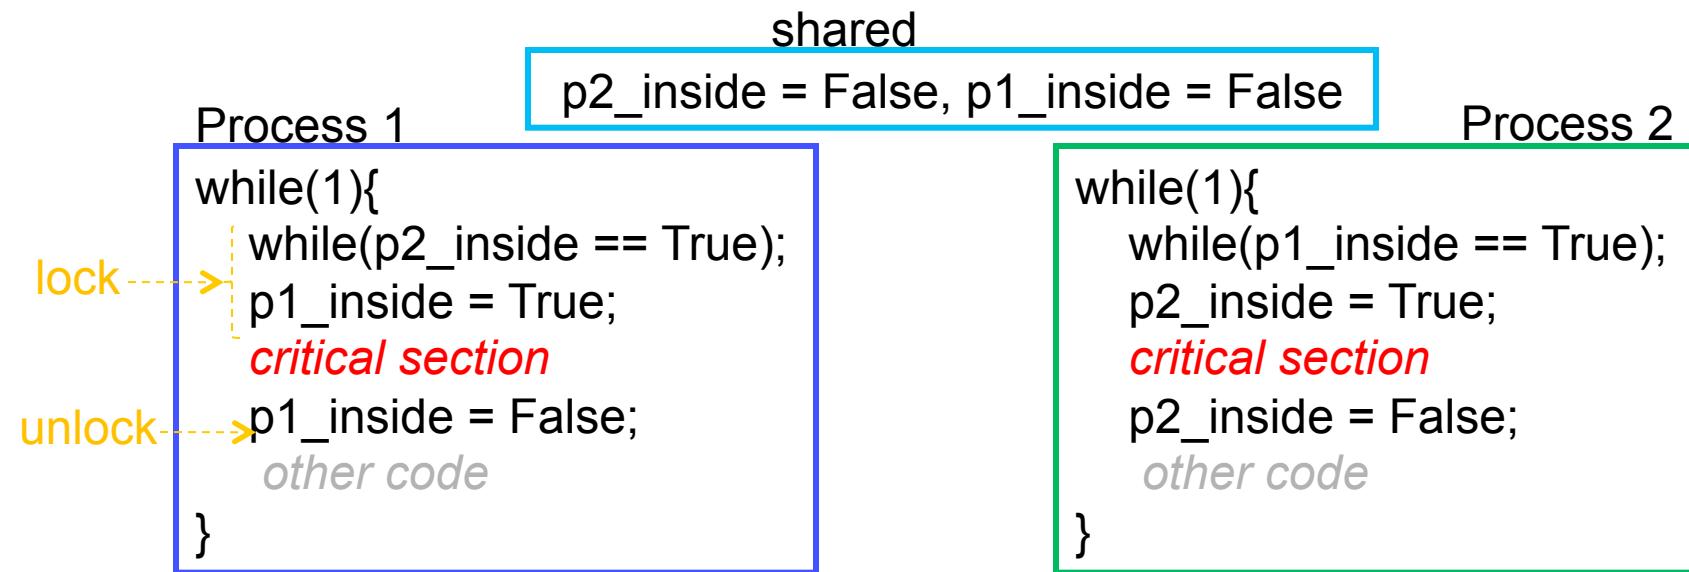
- Achieves mutual exclusion
 - Busy waiting – waste of power and time
 - Needs to alternate execution in critical section
- process1 → process2 → process1 → process2*

Problem with Attempt 1



- Had a common turn flag that was modified by both processes
- This required processes to alternate.
- Possible Solution : Have two flags – one for each process

Software Solution (Attempt 2)



- Need not alternate execution in critical section
- Does not guarantee mutual exclusion

mutex is not guaranteed if both the processes context switched right after exiting the waiting while loop

Attempt 2: No mutual exclusion



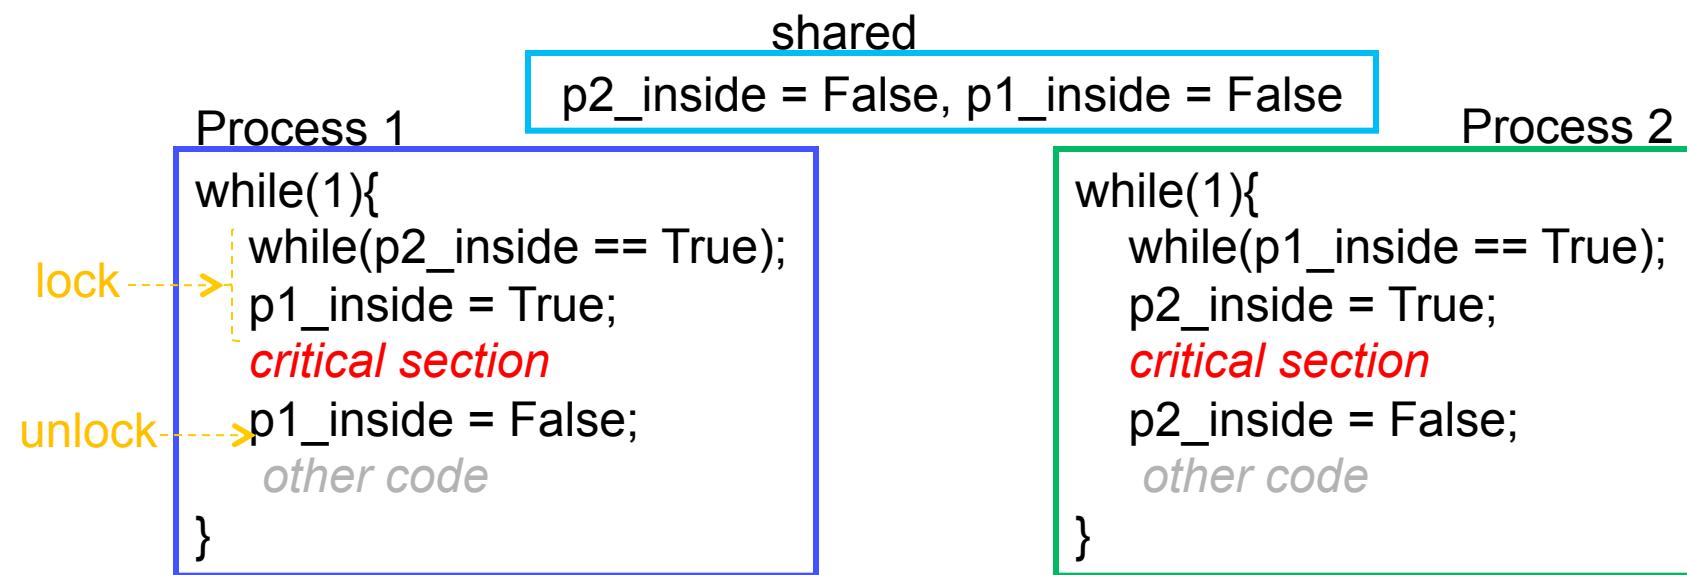
CPU	p1_inside	p2_inside
<code>while(p2_inside == True);</code>	False	False
context switch		
<code>while(p1_inside == True);</code>	False	False
<code>p2_inside = True;</code>	False	True
context switch		
<code>p1_inside = True;</code>	True	True

Both p1 and p2 can enter into the critical section at the same time

```
while(1){  
    while(p2_inside == True);  
    p1_inside = True;  
    critical section  
    p1_inside = False;  
    other code  
}
```

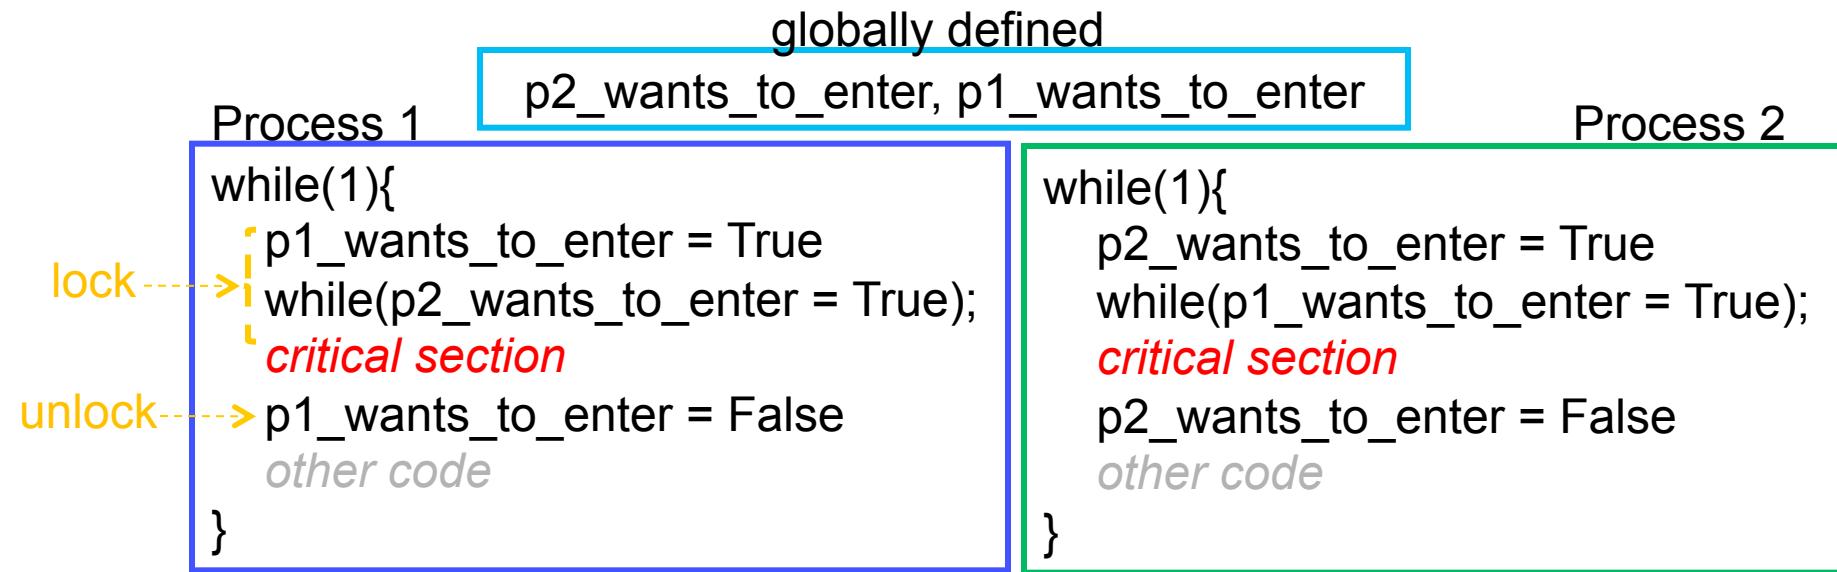
```
while(1){  
    while(p1_inside == True);  
    p2_inside = True;  
    critical section  
    p2_inside = False;  
    other code  
}
```

Problem with Attempt 2



- The flag (`p1_inside`, `p2_inside`), is set after we break from the while loop.

Software Solution (Attempt 3)



- Achieves mutual exclusion
- Does not achieve progress (could deadlock)

Attempt 3: No Progress

↓
time

CPU	p1_inside	p2_inside
p1_wants_to_enter = True	False	False
context switch		
p2_wants_to_enter = True	False	False

There is a tie!!!

Both p1 and p2 will loop infinitely

Progress not achieved

Each process is waiting for the other
this is a deadlock

```
while(1){  
    p2_wants_to_enter = True  
    while(p1_wants_to_enter = True);  
        critical section  
    p2_wants_to_enter = False  
        other code  
}
```

Deadlock

↓
time

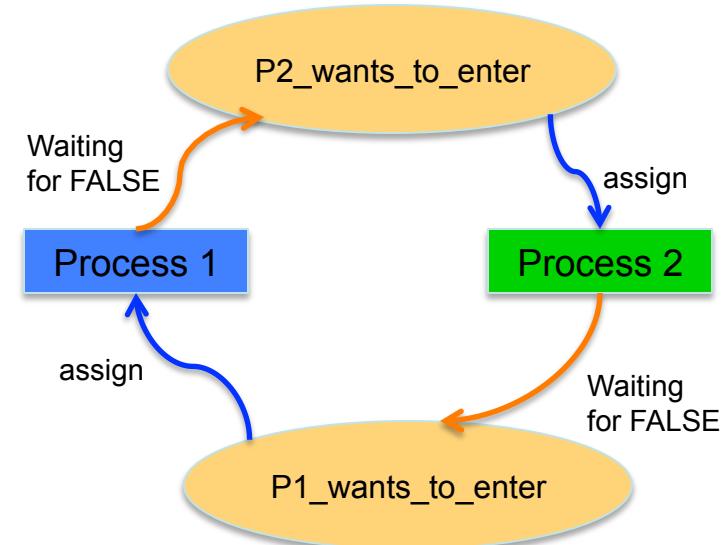
CPU	p1_inside	p2_inside
p1_wants_to_enter = True	False	False
context switch		
p2_wants_to_enter = True	False	False

There is a tie!!!

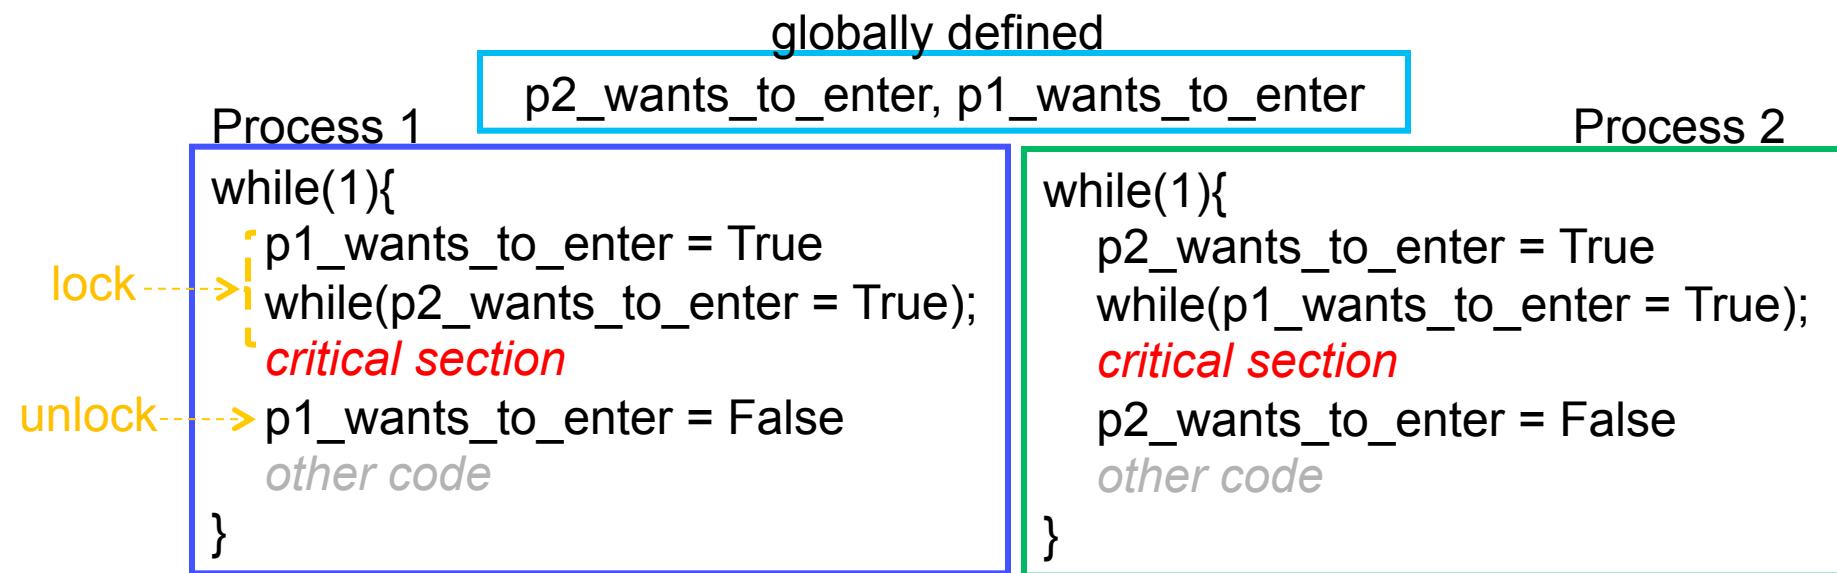
Both p1 and p2 will loop infinitely

Progress not achieved

Each process is waiting for the other
this is a deadlock

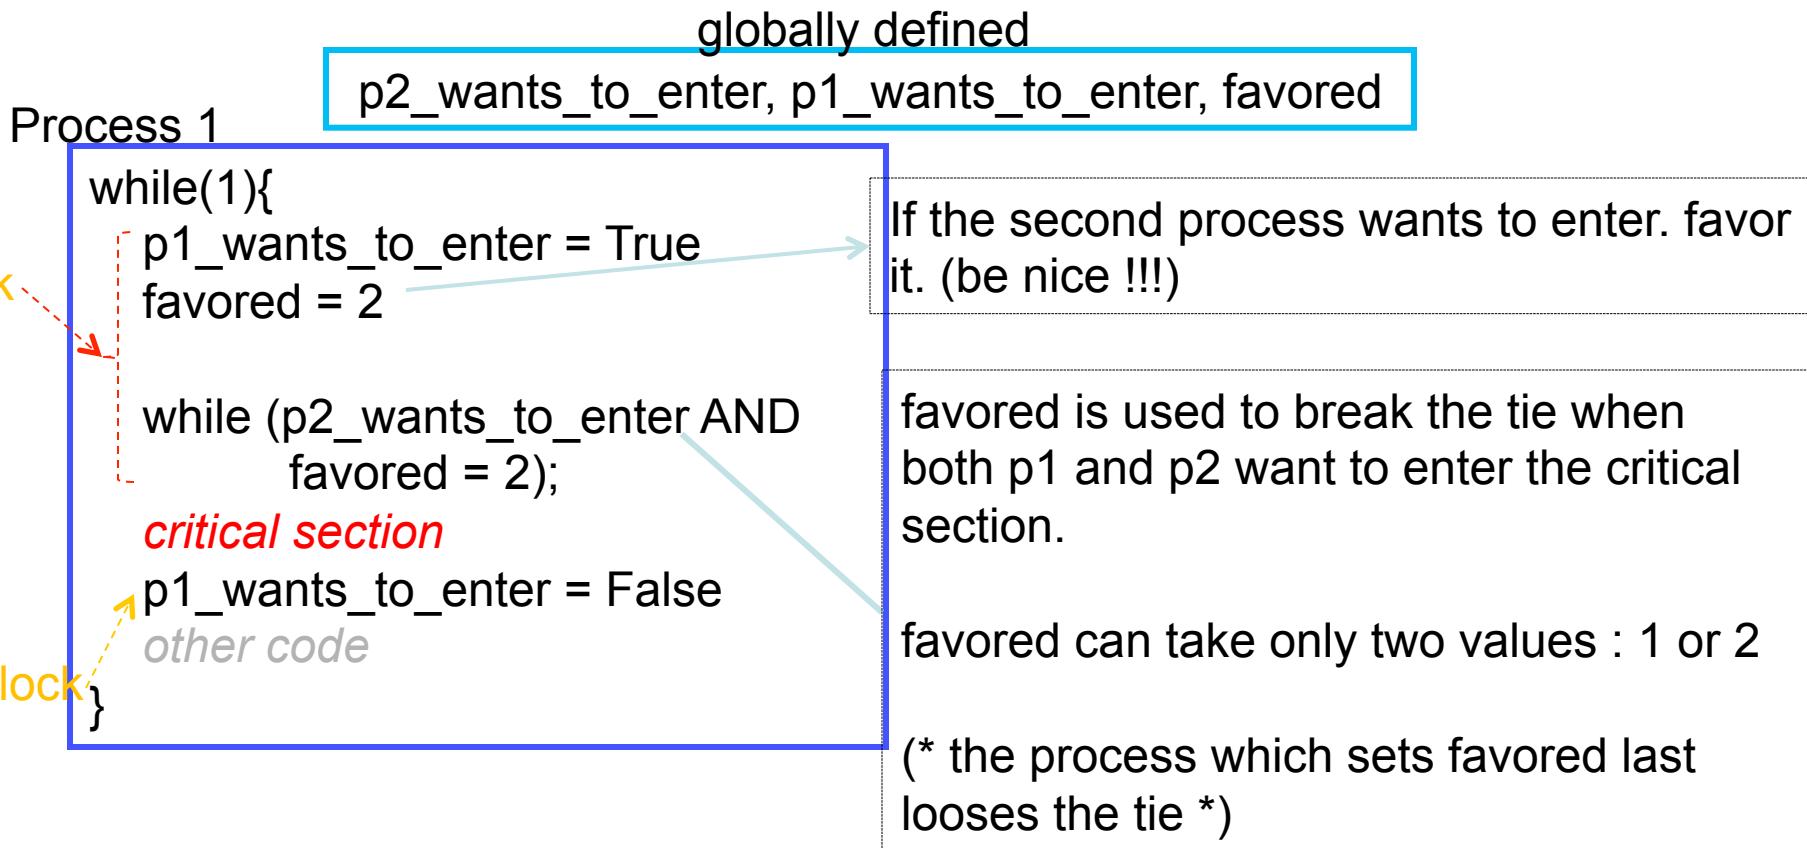


Problem with Attempt 3



- Deadlock
Have a way to break the deadlock

Peterson's Solution



Break the deadlock with a ‘favored’ process

Peterson's Solution

Process 1

```
globally defined  
p2_wants_to_enter, p1_wants_to_enter, favored  
while(1){  
    p1_wants_to_enter = True  
    favored = 2  
  
    while (p2_wants_to_enter AND  
          favored = 2);  
    critical section  
    p1_wants_to_enter = False  
    other code  
}
```

Process 2

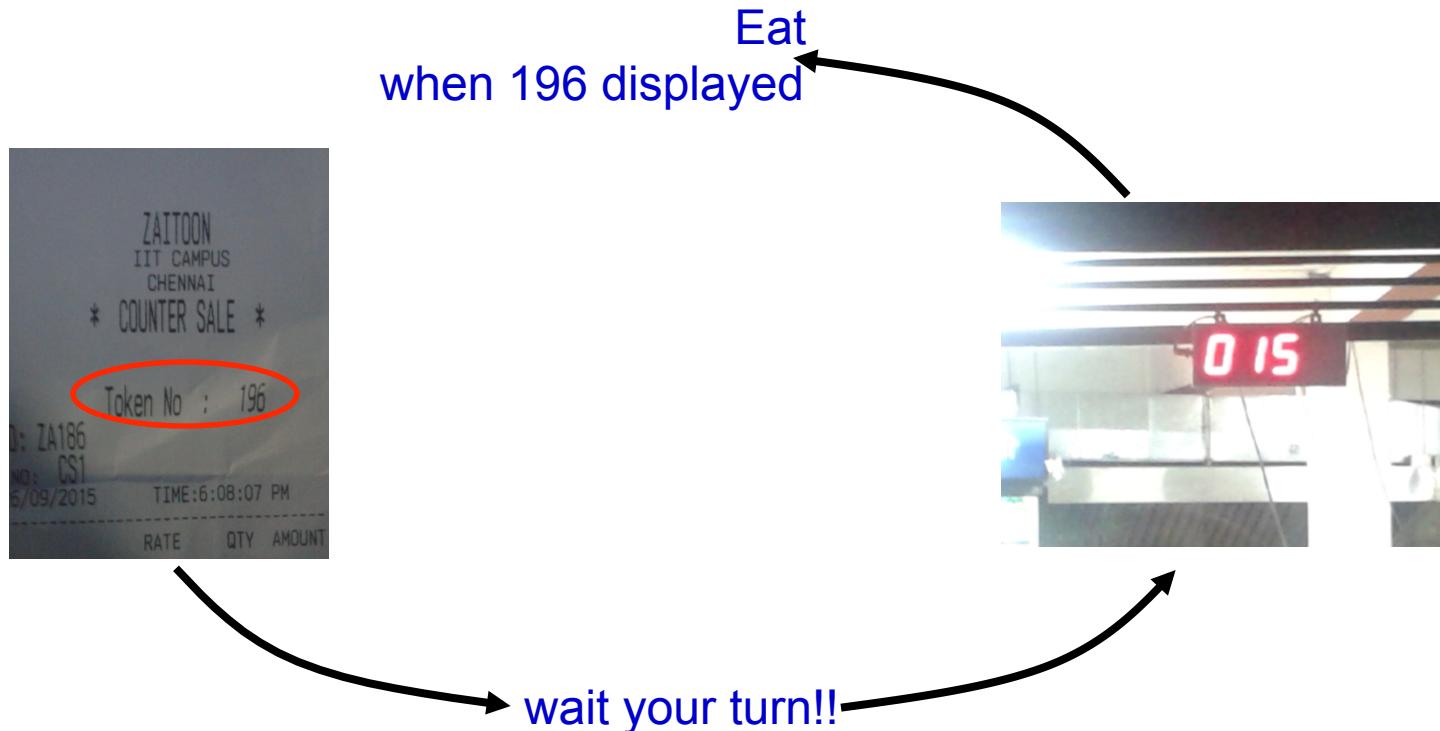
```
while(1){  
    p2_wants_to_enter = True  
    favored = 1  
  
    while (p1_wants_to_enter AND  
          favored = 1);  
    critical section  
    p2_wants_to_enter = False  
    other code  
}
```

- Deadlock broken because favored can only be 1 or 2.
 - Therefore, tie is broken. Only one process will enter the critical section
- Solves Critical Section problem for two processes

this solution is identical to previous attempt and solves mutual exclusion. progress is also guaranteed as when only one process is about to run, it can freely enter the critical section. the bounded wait issue is still not explicitly addressed. this solution is also applicable only for a system of two processes/threads.

Bakery Algorithm

- Synchronization between $N > 2$ processes
- By Leslie Lamport



Simplified Bakery Algorithm

- Processes numbered 0 to N-1
- num is an array N integers (initially 0).
 - Each entry corresponds to a process

```
lock(i){    this is basically equivalent to getting a token that is one greater than the highest value token currently out there
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1
    for(p = 0; p < N; ++p){
        while (num[p] != 0 and num[p] < num[i]);
    }    you wait until all the other who have a token with number less than you is served
}
```

critical section

```
unlock(i){    when service is done, forfeit the token
    num[i] = 0;
}
```

This is at the doorway!!!
It has to be atomic
to ensure two processes
do not get the same token

this is very important. if two processes get the same token number, both of them can enter the critical section at the same time and mutex is violated.

Simplified Bakery Algorithm (example)

- Processes numbered 0 to N-1
- num is an array N integers (initially 0).
 - Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

critical section

	P1	P2	P3	P4	P5
0	0	0	0	2	3

```
unlock(i){  
    num[i] = 0;  
}
```

Simplified Bakery Algorithm (example)

Processes numbered 0 to N-1

num is an array N integers (initially 0).

Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

	P1	P2	P3	P4	P5
	0	0	0	0	0

Simplified Bakery Algorithm (why atomic doorway?)

- Processes numbered 0 to N-1
- num is an array N integers (initially 0).
 - Each entry corresponds to a process

```
lock(i){  
    num[i] = MAX(num[0], num[1], ...., num[N-1]) + 1  
    for(p = 0; p < N; ++p){  
        while (num[p] != 0 and num[p] < num[i]);  
    }  
}
```

This is at the doorway!!!
Assume it is not atomic

critical section

P1	P2	P3	P4	P5
0	0	0	0	0

```
unlock(i){  
    num[i] = 0;  
}
```

P4 and P5 can enter the critical section at the same time.

Original Bakery Algorithm (making MAX atomic)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

doorway

Choosing ensures that a process
Is not at the doorway
i.e., the process is not 'choosing'
a value for num

$(a, b) < (c, d)$ which is equivalent to: $(a < c)$ or $((a == c) \text{ and } (b < d))$

Original Bakery Algorithm (making MAX atomic)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

we basically put a mutex around the token allocation process itself
also there is a modification to allow breaking of ties when there is a deadlock

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

Favor one process when there is a conflict.

If there are two processes, with the same num value, favor the process with the smaller id (i)

$(a, b) < (c, d)$ which is equivalent to: $(a < c)$ or $((a == c) \text{ and } (b < d))$

Original Bakery Algorithm (example)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

	P1	P2	P3	P4	P5
	0	0	0	0	0

$(a, b) < (c, d)$ which is equivalent to: $(a < c)$ or $((a == c) \text{ and } (b < d))$

Original Bakery Algorithm (example)

- Without atomic operation assumptions
- Introduce an array of N Booleans: *choosing*, initially all values False.

```
lock(i){  
    choosing[i] = True  
    num[i] = MAX(num[0], num[1], ..., num[N-1]) + 1  
    choosing[i] = False  
    for(p = 0; p < N; ++p){  
        while (choosing[p]);  
        while (num[p] != 0 and (num[p],p)<(num[i],i));  
    }  
}
```

doorway

critical section

```
unlock(i){  
    num[i] = 0;  
}
```

P1	P2	P3	P4	P5
0	3	0	2	2

$(a, b) < (c, d)$ which is equivalent to: $(a < c)$ or $((a == c) \text{ and } (b < d))$

How to Implement Locking (Hardware Solutions and Usage)

Analyze this

- Does this scheme provide mutual exclusion?

Process 1

```
while(1){  
    while(lock != 0);  
    lock= 1; // lock  
    critical section  
    lock = 0; // unlock  
    other code  
}
```

lock=0

Process 2

```
while(1){  
    while(lock != 0);  
    lock = 1; // lock  
    critical section  
    lock = 0; // unlock  
    other code  
}
```

No

```
lock = 0  
P1: while(lock != 0);  
P2: while(lock != 0);  
P2: lock = 1;  
P1: lock = 1;  
.... Both processes in critical section
```

context switch

If only...

- We could make this operation atomic

Process 1

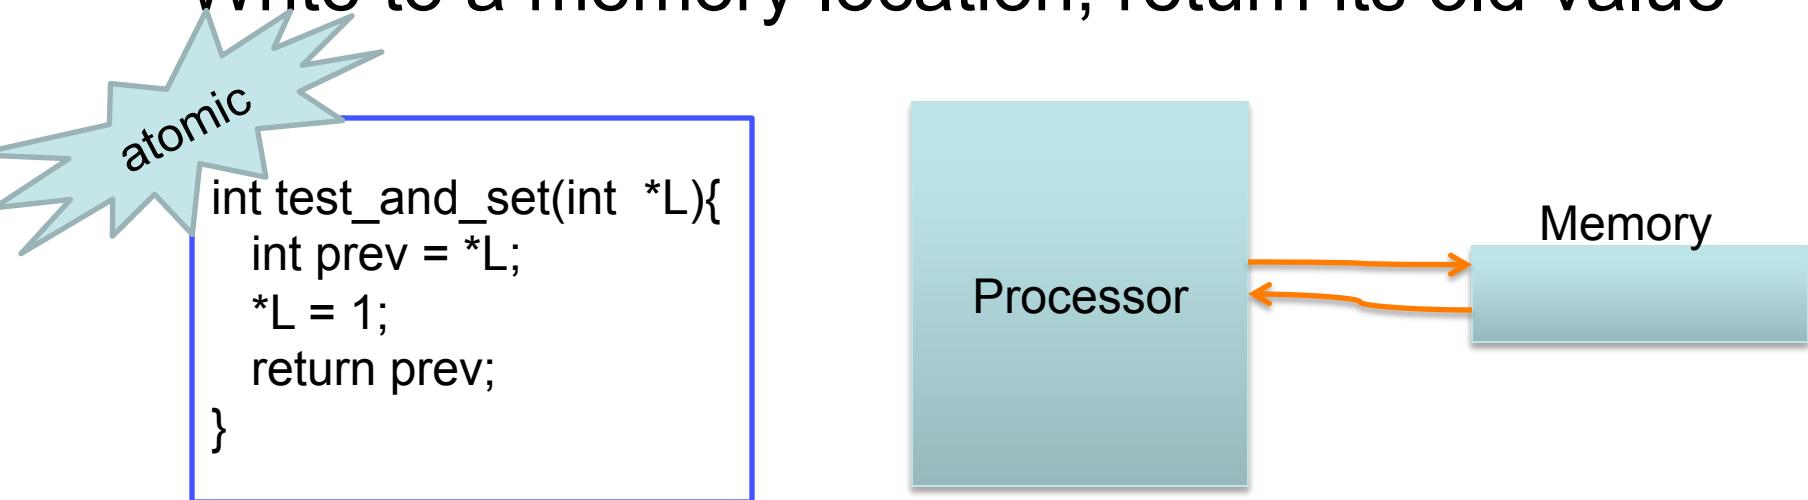
```
while(1){  
    while(lock != 0);  
    lock= 1; // lock  
critical section  
    lock = 0; // unlock  
    other code  
}
```

Make atomic

Hardware to the rescue....

Hardware Support (Test & Set Instruction)

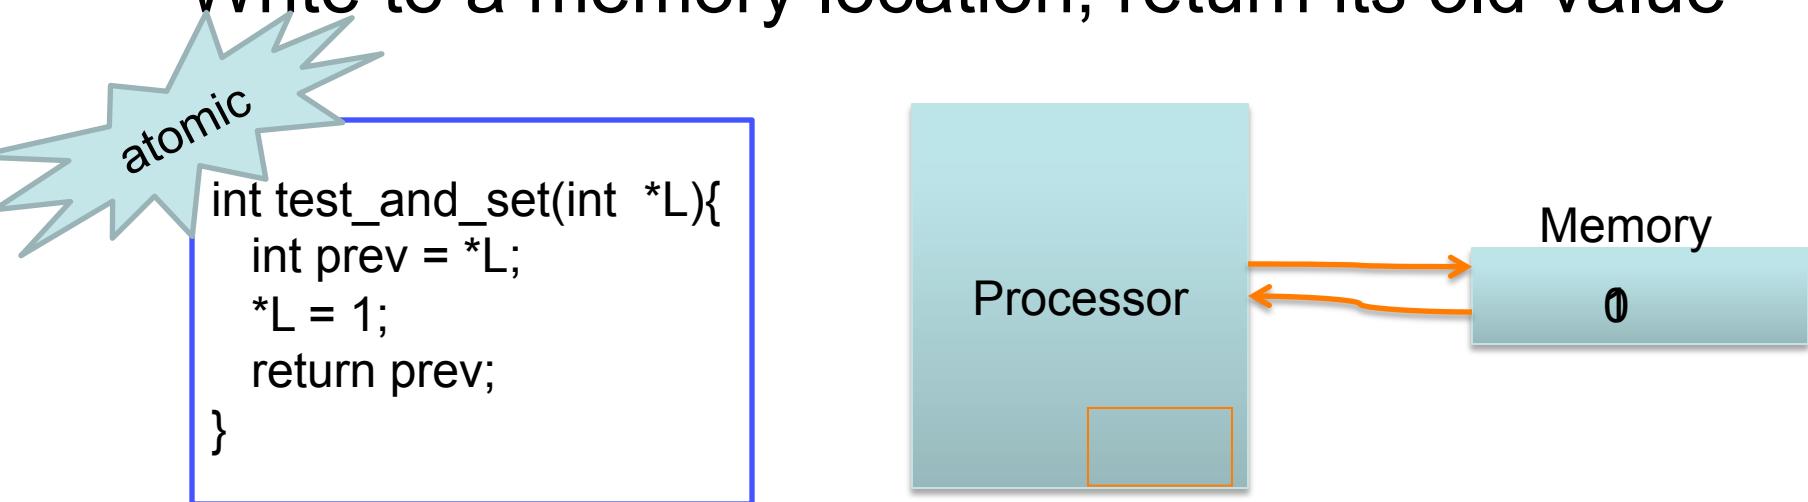
- Write to a memory location, return its old value



equivalent software representation
(the entire function is executed atomically)

Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



equivalent software representation
(the entire function is executed atomically)

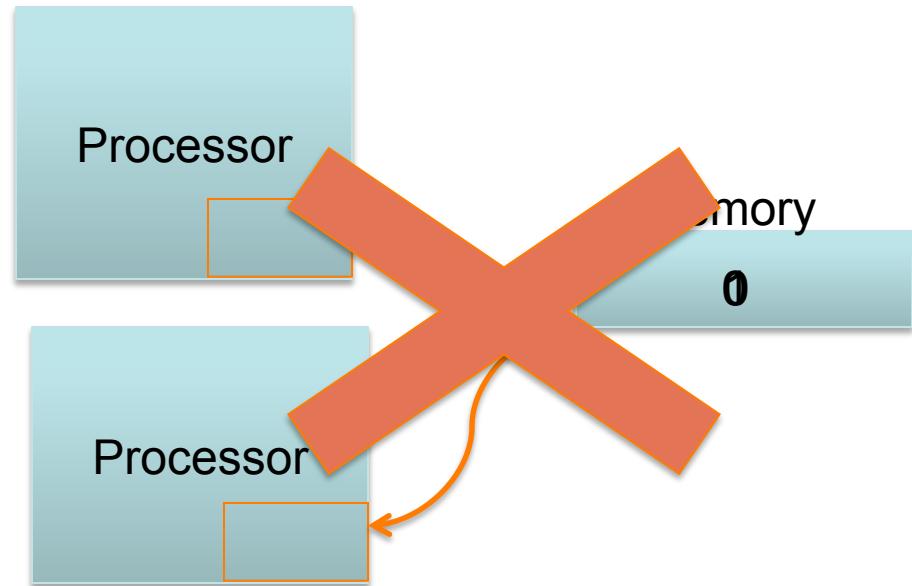
Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value

atomic

```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

equivalent software representation
(the entire function is executed
atomically)

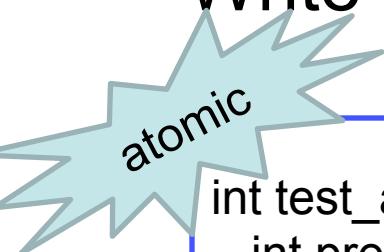


this is weird - when multiple cpus attempt test_set at the same time, there can be significant stalling

Why does this work? If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

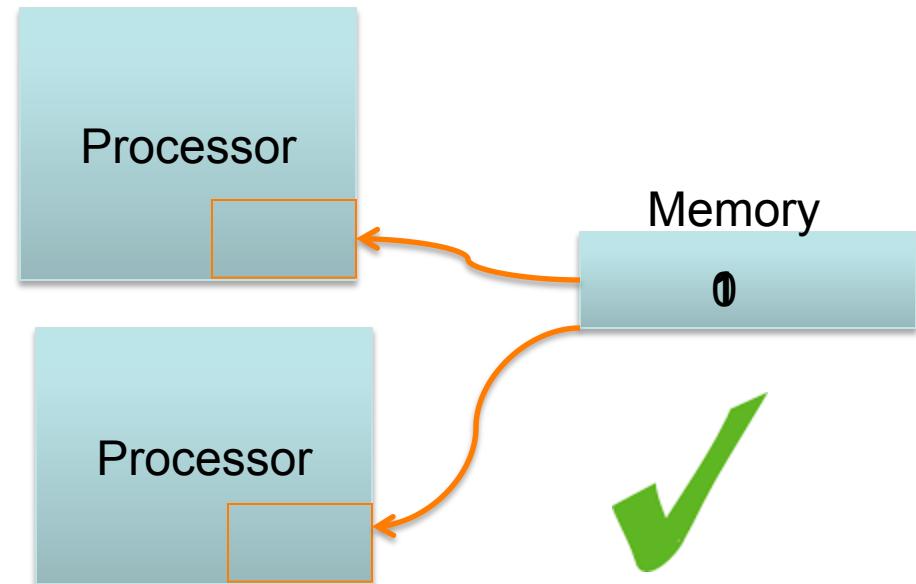
Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

equivalent software representation
(the entire function is executed
atomically)



Why does this work? If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

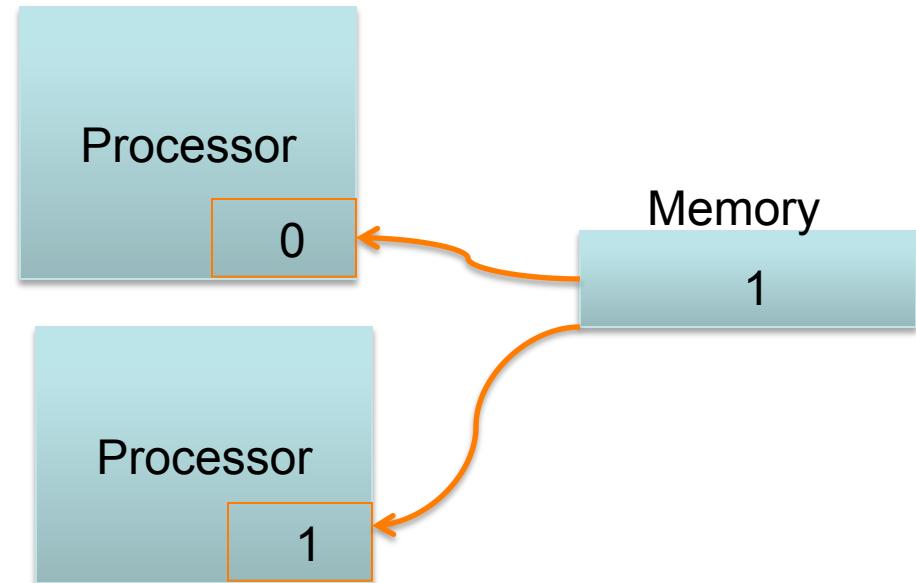
Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value

atomic

```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

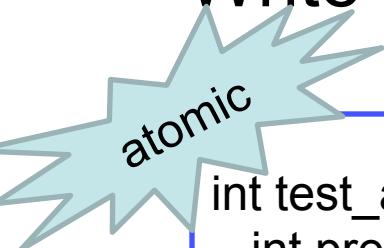
equivalent software representation
(the entire function is executed
atomically)



Why does this work? If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

Hardware Support (Test & Set Instruction)

- Write to a memory location, return its old value



atomic

```
int test_and_set(int *L){  
    int prev = *L;  
    *L = 1;  
    return prev;  
}
```

```
while(1){  
    while(test_and_set(&lock) == 1);  
    critical section  
    lock = 0; // unlock  
    other code  
}
```

equivalent software representation
(the entire function is executed atomically)

Usage for locking

Why does this work? If two CPUs execute `test_and_set` at the same time, the hardware ensures that one `test_and_set` does both its steps before the other one starts.

So the first invocation of `test_and_set` will read a 0 and set lock to 1 and return. The second `test_and_set` invocation will then see lock as 1, and will loop continuously until lock becomes 0

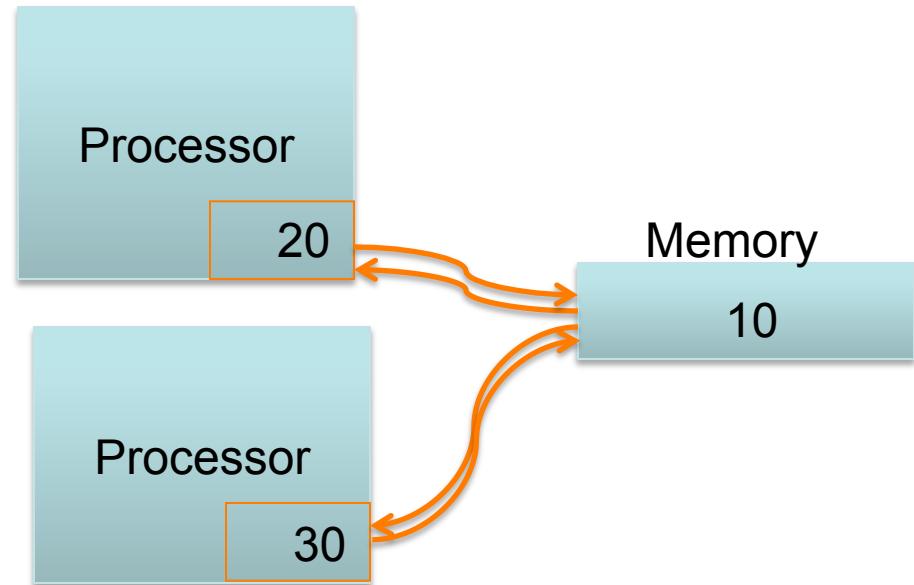
Intel Hardware Support (xchg Instruction)

- Write to a memory location, return its old value



```
int xchg(int *L, int v){  
    int prev = *L;  
    *L = v;  
    return prev;  
}
```

equivalent software representation
(the entire function is executed
atomically)

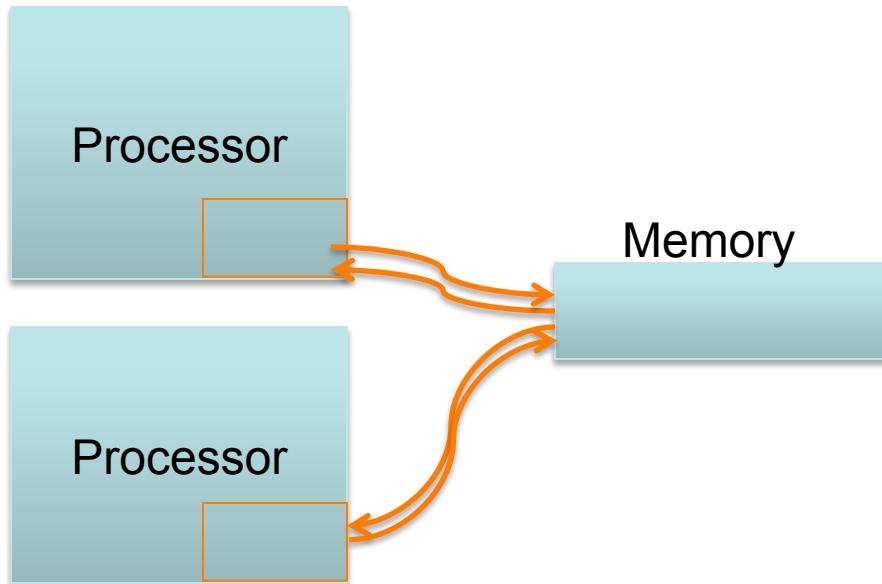


Why does this work? If two CPUs execute xchg at the same time, the hardware ensures that one xchg completes, only then the second xchg starts.

Intel Hardware Support (using xchg instruction)

Note. %eax is returned
typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

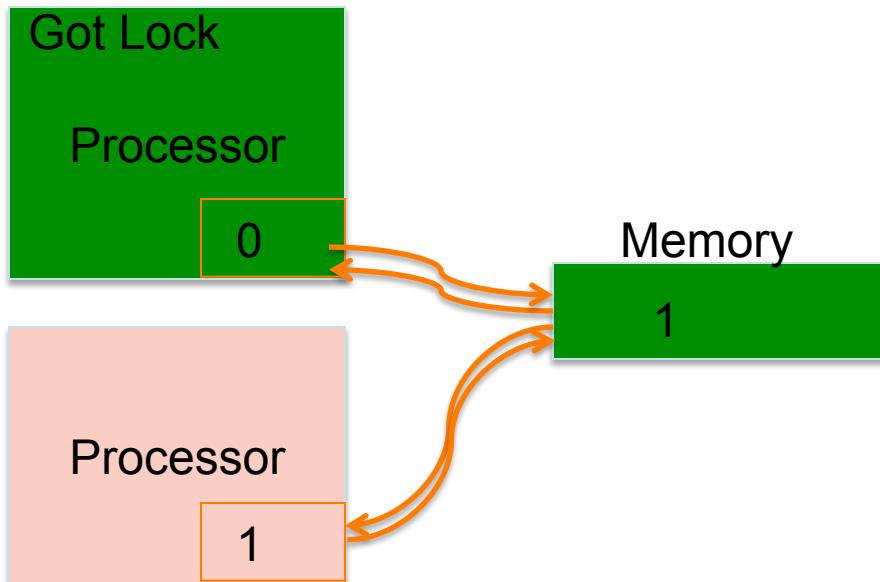
```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```

Intel Hardware Support (using xchg instruction)

Note. %eax is returned
typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

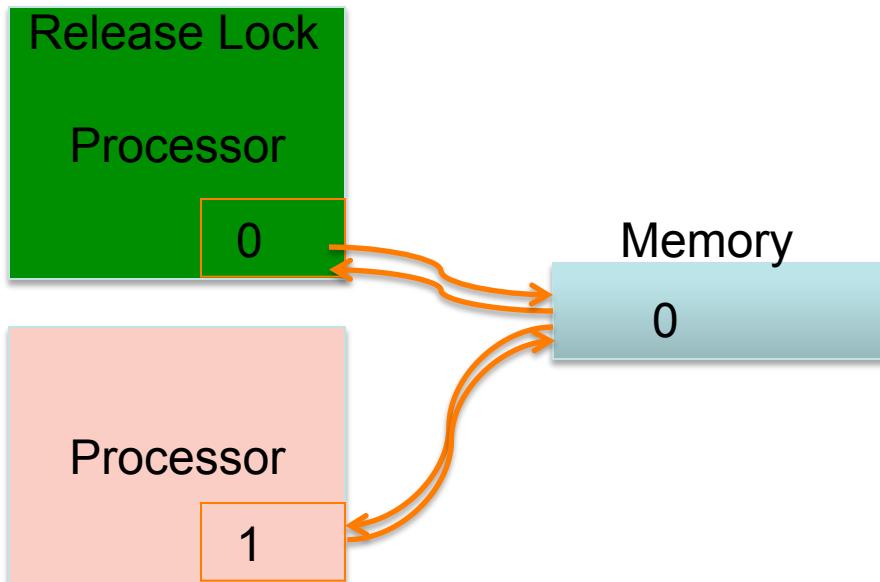
```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```

Intel Hardware Support (using xchg instruction)

Note. %eax is returned
typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

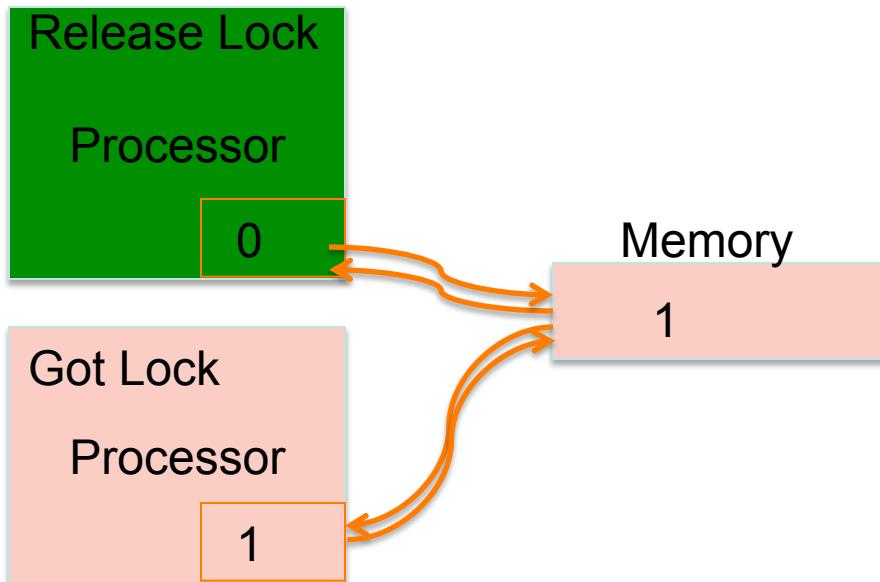
```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```

Intel Hardware Support (using xchg instruction)

Note. %eax is returned
typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

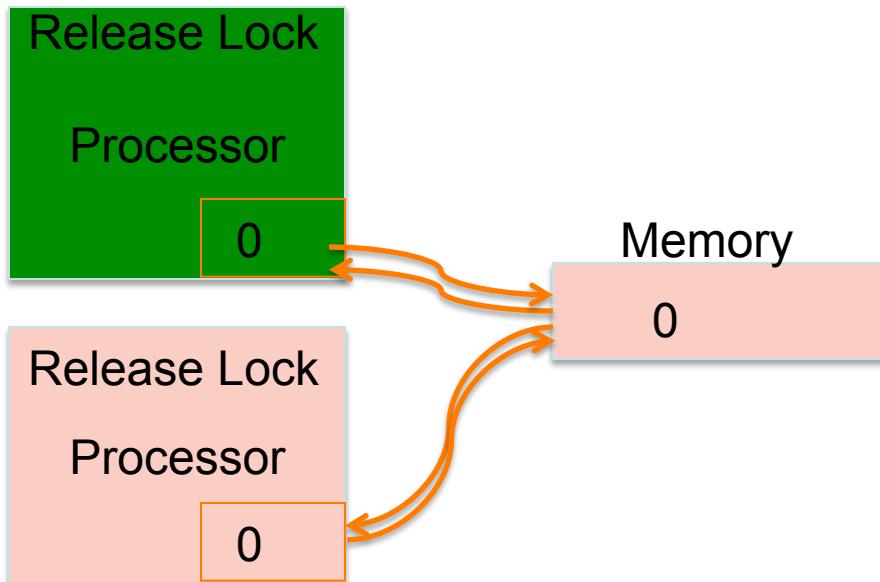
```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```

Intel Hardware Support (using xchg instruction)

Note. %eax is returned
typical usage :

xchg reg, mem



```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

```
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}
```

```
void release(int *locked){  
    locked = 0;  
}
```

High Level Constructs

- Spinlock
- Mutex
- Semaphore

Spinlocks Usage

Process 1

```
acquire(&locked)  
critical section  
release(&locked)
```

Process 2

```
acquire(&locked)  
critical section  
release(&locked)
```

- One process will **acquire** the lock
- The other will wait in a loop repeatedly checking if the lock is available
- The lock becomes available when the former process **releases** it

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void acquire(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
    }  
}  
  
void release(int *locked){  
    locked = 0;  
}
```

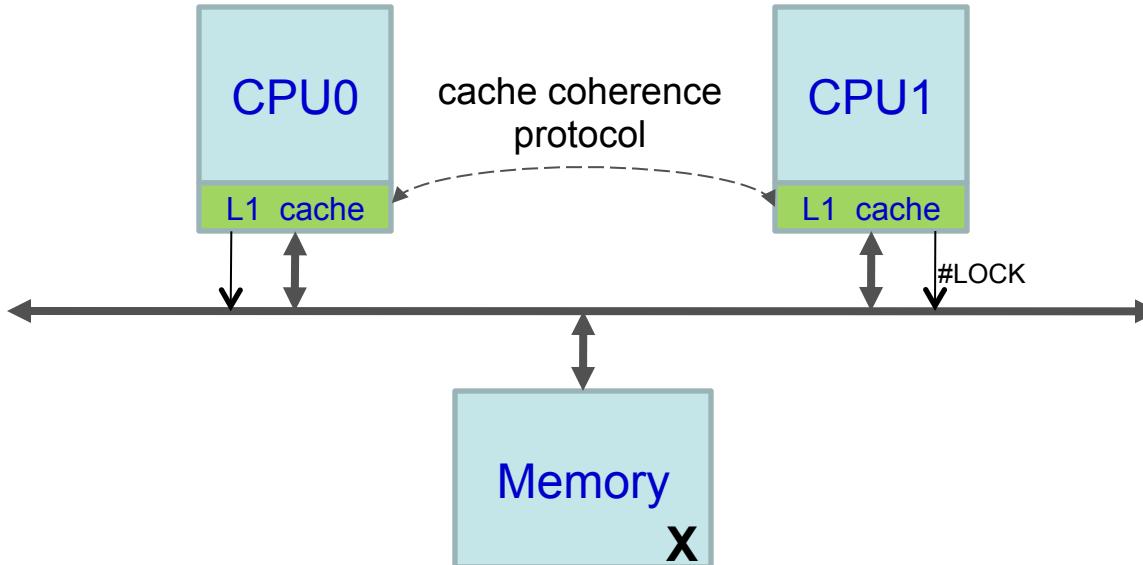
Issues with Spinlocks

`xchg %eax, X`

- No compiler optimizations should be allowed
 - Should not make X a register variable
 - Write the loop in assembly or use volatile
- Should not reorder **memory** loads and stores
 - Use serialized instructions (which forces instructions not to be reordered)
 - Luckily xchg is already implements serialization

More issues with Spinlocks

xchg %eax, X



- No caching of (X) possible. All xchg operations are bus transactions.
 - CPU asserts the LOCK, to inform that there is a 'locked' memory access
- acquire function in spinlock invokes xchg in a loop...each operation is a bus transaction **huge performance hits**

this bus transaction would block the memory and not permit other processes from accessing the memory close to the lock

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}
```

A better acquire

```
void acquire(int *locked){  
    reg = 1  
    while(1)  
        if(xchg(locked, reg) == 0)  
            break;  
}
```

```
void acquire(int *locked) {  
    reg = 1;  
    while (xchg(locked, reg) == 1)  
        while (*locked == 1);  
}
```



Original.

Loop with xchg.
Bus transactions.
Huge overheads

Better way

Outer loop changes the value of locked
inner loop only reads the value of
locked. This allows caching of
locked.

Access cache instead of memory.

this would then require cache coherence to allow snooping
so that when another process releases the lock, the new one can get
the lock.

Spinlocks (when should it be used?)

- Characteristic : busy waiting
 - Useful for short critical sections, where much CPU time is not wasted waiting
 - eg. To increment a counter, access an array element, etc.
 - Not useful, when the period of wait is unpredictable or will take a long time
 - eg. Not good to read page from disk.
 - Use mutex instead (...mutex)

Spinlock in pthreads

```
#include <pthread.h>
#include <stdio.h>

int global_counter;
pthread_spinlock_t splk;

void *thread_fn(void *arg){
    long id = (long) arg;
    while(1){
        pthread_spin_lock(&splk); → lock
        if (id == 1) global_counter++;
        else global_counter--;
        pthread_spin_unlock(&splk); → unlock
        printf("%d(%d)\n", id, global_counter);
        sleep(1);
    }

    return NULL;
}

int main(){
    pthread_t t1, t2;

    pthread_spin_init(&splk, PTHREAD_PROCESS_PRIVATE); → create spinlock
    pthread_create(&t1, NULL, thread_fn, (void *)1);
    pthread_create(&t2, NULL, thread_fn, (void *)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    pthread_spin_destroy(&splk); → destroy spinlock
    printf("Exiting main\n");
    return 0;
}
```

Mutexes

- Can we do better than busy waiting?
 - If critical section is locked then yield CPU
 - Go to a SLEEP state
 - While unlocking, wake up sleeping process

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void lock(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else  
            sleep();  
    }  
}  
  
void unlock(int *locked){  
    locked = 0;  
    wakeup();  
}
```

Thundering Herd Problem

- A large number of processes wake up (almost simultaneously) when the event occurs.
 - All waiting processes wake up
 - Leading to several context switches
 - All processes go back to sleep except for one, which gets the critical section
 - Large number of context switches
 - Could lead to starvation

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void lock(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else  
            sleep();  
    }  
}  
  
void unlock(int *locked){  
    locked = 0;  
    wakeup();  
}
```

these contextswitches are useless as nearly all process go back to sleep

this can be fixed by adding all waiting processes to a FIFO like queue where during wakeup only the front of the queue is woken up
this would then drastically decrease the number of useless context switches

Thundering Herd Problem

- The Solution
 - When entering critical section, push into a queue before blocking
 - When exiting critical section, wake up only the first process in the queue

```
int xchg(addr, value){  
    %eax = value  
    xchg %eax, (addr)  
}  
  
void lock(int *locked){  
    while(1){  
        if(xchg(locked, 1) == 0)  
            break;  
        else{  
            // add this process to Queue  
            sleep();  
        }  
    }  
}  
  
void unlock(int *locked){  
    locked = 0;  
    // remove process P from queue  
    wakeup(P)  
}
```

pthread Mutex

- `pthread_mutex_lock`
- `pthread_mutex_unlock`

Locks and Priorities

- What happens when a high priority task requests a lock, while a low priority task is in the critical section
 - Priority Inversion
 - Possible solution
 - Priority Inheritance

whenever a low priority task holds a lock which a high priority task might also access,
the low priority task is temporarily boosted to a higher priority level after which the high priority task can
interrupt the low priority process and perform its operations

Interesting Read : Mass Pathfinder

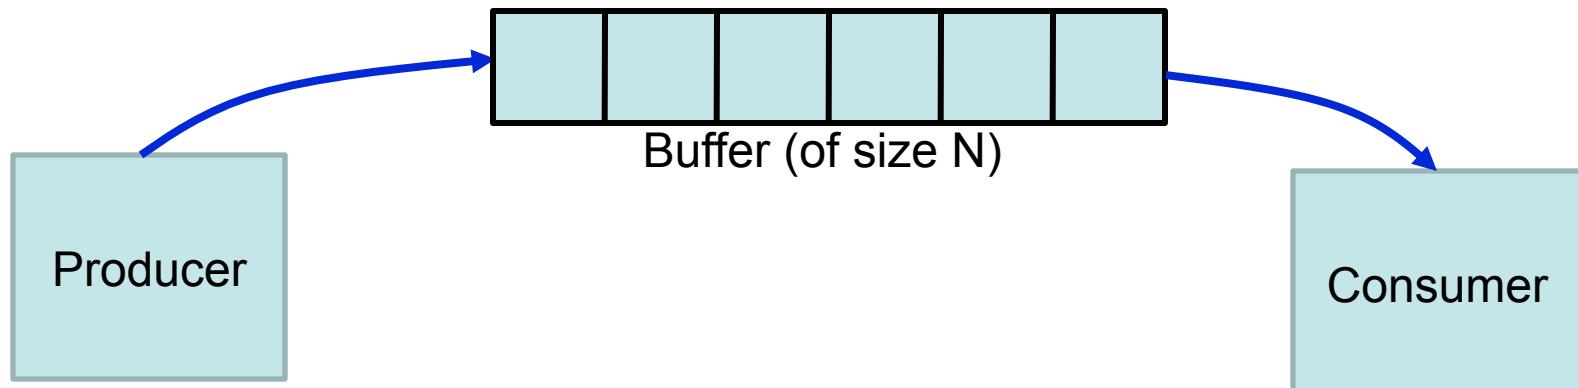
http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/mars_pathfinder.html

Semaphores

Producer – Consumer Problems

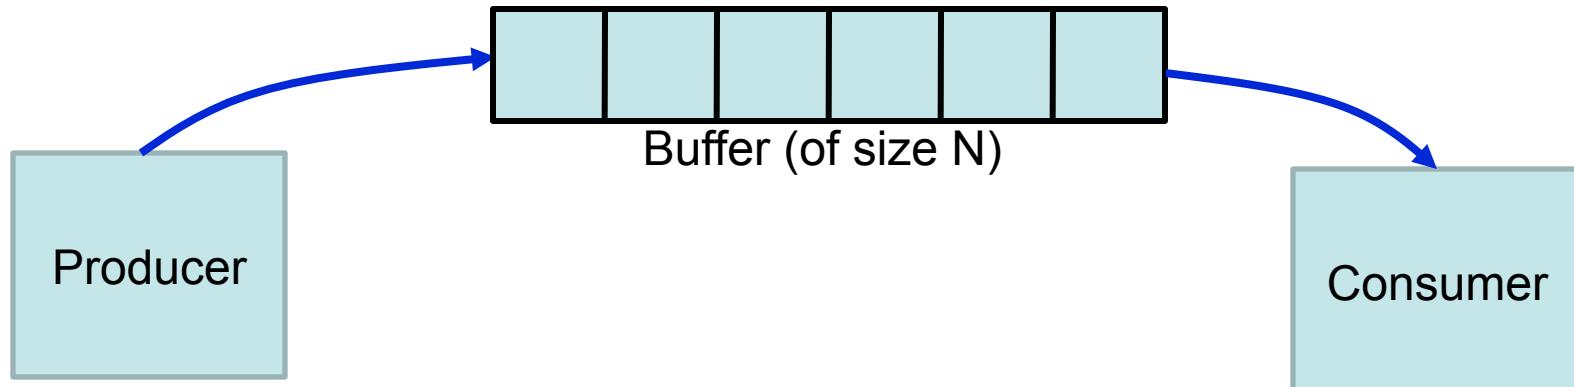
- Also known as *Bounded buffer Problem*
- Producer produces and stores in buffer, Consumer consumes from buffer

the whole problem arises when the size of the buffer is finite



Producer – Consumer Problems

- Also known as *Bounded buffer Problem*
- Producer produces and stores in buffer, Consumer consumes from buffer
- Trouble when
 - Producer produces, but buffer is full
 - Consumer consumes, but buffer is empty



Producer-Consumer Code

Buffer of size N

```
int count=0;  
Mutex mutex, empty, full;
```

```
1 void producer(){  
2     while(TRUE){  
3         item = produce_item();  
4         if (count == N) sleep(empty);  
5         lock(mutex);  
6         insert_item(item); // into buffer  
7         count++;  
8         unlock(mutex);  
9         if (count == 1) wakeup(full);  
10    }  
}
```

```
1 void consumer(){  
2     while(TRUE){  
3         if (count == 0) sleep(full);  
4         lock(mutex);  
5         item = remove_item(); // from buffer  
6         count--;  
7         unlock(mutex);  
8         if (count == N-1) wakeup(empty);  
9         consume_item(item);  
10    }  
}
```

the whole point of this construct is to handle the edge cases where the buffer is FULL or EMPTY

note that the sleeps and wakeup are done on separate mutexes

Producer-Consumer Code

Buffer of size N

int count=0;

Mutex mutex, empty, full;

```
1 void producer(){
2     while(TRUE){
3         item = produce_item();
4         if (count == N) sleep(empty);
5         lock(mutex);
6         insert_item(item); // ...
7         count++;
8         unlock(mutex);
9         if (count == 1) wakeup(full);
10    }
```

Read count value
Test count = 0

```
1 void consumer(){
2     while(TRUE){
3         if (count == 0) sleep(full);
4         lock(mutex);
5         item = remove_item(); // from buffer
6         count--;
7         unlock(mutex);
8         if (count == N-1) wakeup(empty);
9         consume_item(item);
10    }
```

Lost Wakeups

- Consider the following context of instructions
- Assume buffer is initially empty

context switch

```
3 read count value // count ← 0
3 item = produce_item();
5 lock(mutex);
6 insert_item(item); // into buffer
7 count++; // count = 1
8 unlock(mutex)
9 test (count == 1) // yes
9 signal(full);
3 test (count == 0) // yes
3 wait();
```

Note, the wakeup is lost.
Consumer waits even though buffer is not empty.
Eventually producer and consumer will wait infinitely

consumer
still uses the old value of count (ie 0)

Semaphores

- Proposed by Dijkstra in 1965
- Functions **down** and **up** must be atomic
- **down** also called **P** (Proberen Dutch for try)
- **up** also called **V** (Verhogen, Dutch form make higher)
- Can have different variants
 - Such as blocking, non-blocking
- If S is initially set to 1,
 - Blocking semaphore similar to a Mutex
 - Non-blocking semaphore similar to a spinlock

both of these functions are atomic

```
void down(int *S){  
    while( *S <= 0);  
    *S--;  
}  
  
void up(int *S){  
    *S++;  
}
```

in a blocking semaphore would sleep and wakeup in some conditions

Producer-Consumer with Semaphores

Buffer of size N

int count;

```
void producer(){
    while(TRUE){
        item = produce_item();
        down(empty);
        wait(mutex);
        insert_item(item); // into buffer
        signal(mutex);
        up(full);
    }
}
```

full = 0, empty = N

```
void consumer(){
    while(TRUE){
        down(full);
        wait(mutex);
        item = remove_item(); // from buffer
        signal(mutex);
        up(empty);
        consume_item(item);
    }
}
```

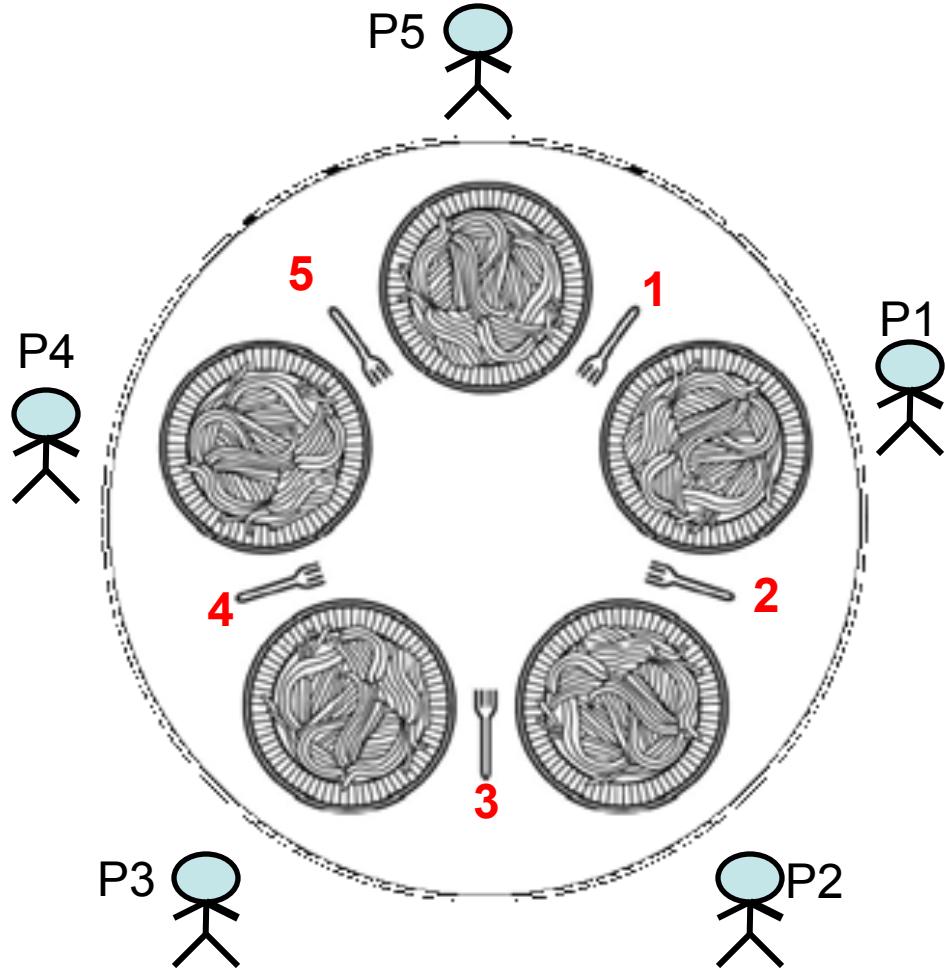
this does not have the problem that was seen when mutexes are used in this place

POSIX semaphores

- `sem_init`
- `sem_wait`
- `sem_post`
- `sem_getvalue`
- `sem_destroy`

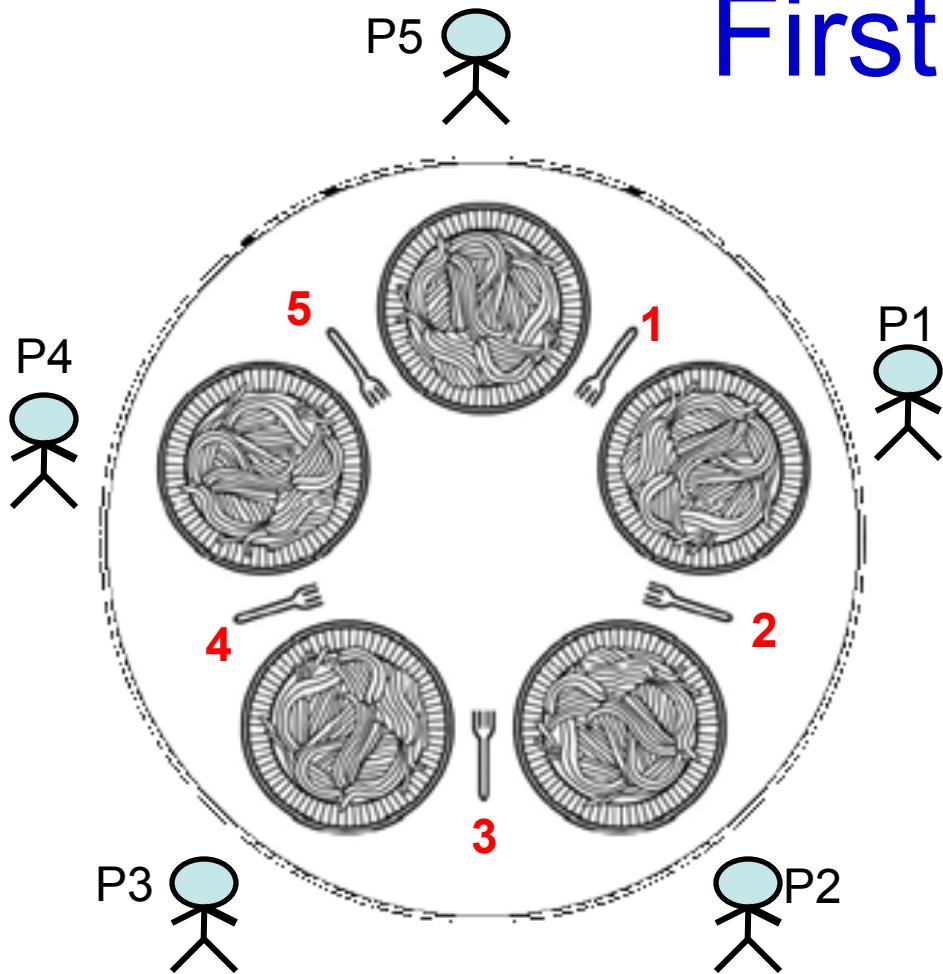
Dining Philosophers Problem

Dining Philosophers Problem



- Philosophers either think or eat
 - To eat, a philosopher needs to hold both forks (the one on his left and the one on his right)
 - If the philosopher is not eating, he is thinking.
-
- **Problem Statement :** Develop an algorithm where no philosopher starves.

First Try

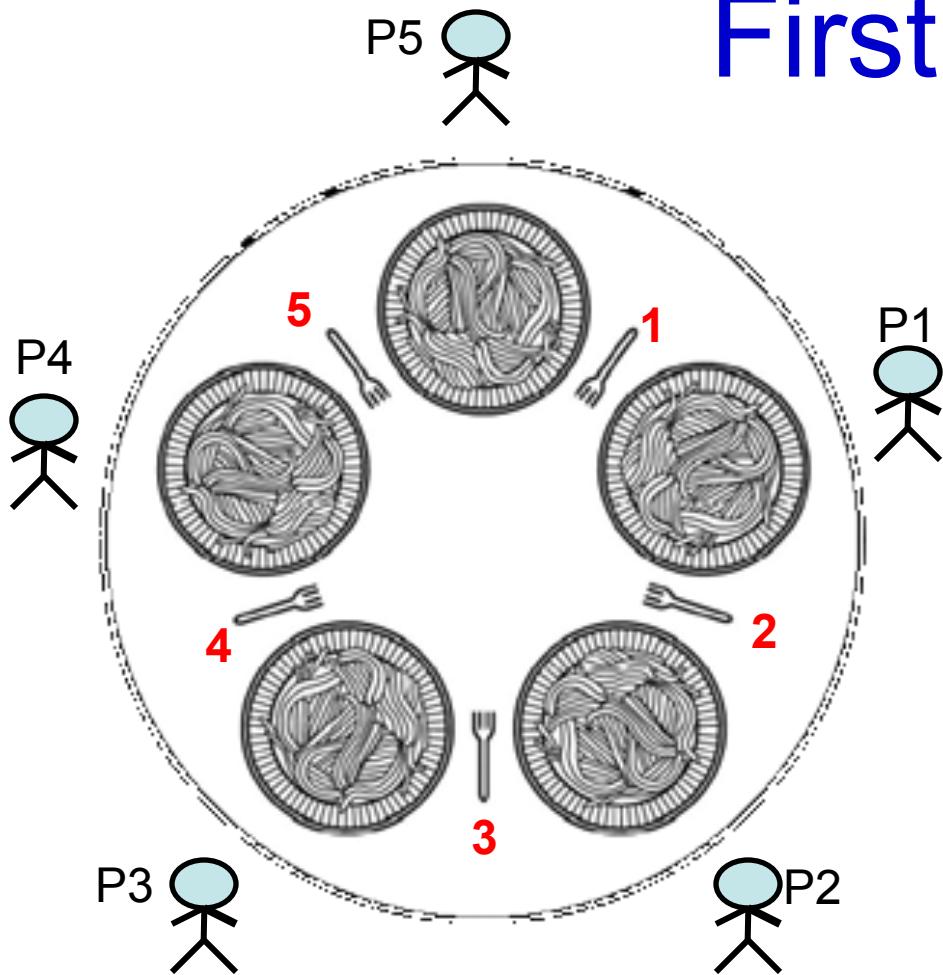


```
#define N 5
int forks = {1,2,3,4,5,1};

void philosopher(int i){
    while(TRUE){
        think(); // for some_time
        take_fork(i);
        take_fork(i + 1);
        eat();
        put_fork(i);
        put_fork(i + 1);
    }
}
```

What happens if only philosophers P1 and P3 are always given the priority?
P2, P4, and P5 starves... so scheme needs to be fair

First Try



```
#define N 5
int forks = {1,2,3,4,5,1};

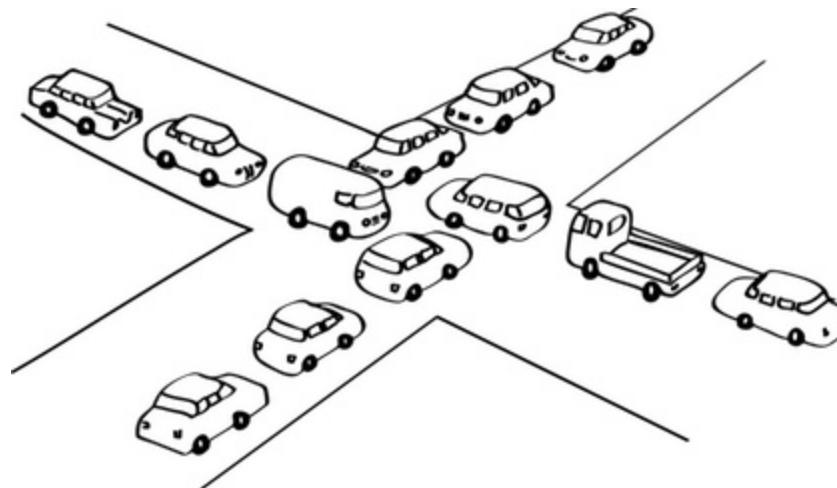
void philosopher(int i){
    while(TRUE){
        think(); // for some_time
        take_fork(i);
        take_fork(i + 1);
        eat();
        put_fork(i);
        put_fork(i + 1);
    }
}
```

What happens if all philosophers decide to pick up their left forks at the same time?
Possible starvation due to deadlock

this is when one guy is waiting for another and this process is in a cycle - with someone finally waiting for someone who initially started

Deadlocks

- A situation where programs continue to run indefinitely without making any progress
- Each program is waiting for an event that another process can cause



Second try

- Take fork i, check if fork (i+1) is available
- Imagine,
 - All philosophers start at the same time
 - Run simultaneously
 - And think for the same time
- This could lead to philosophers taking fork and putting it down continuously.
a deadlock.
- A better alternative
 - Philosophers wait a random time before take_fork(i)
 - Less likelihood of deadlock.
 - Used in schemes such as Ethernet

```
#define N 5
int forks = {1,2,3,4,5,1};

void philosopher(int i){
    while(TRUE){
        think();
        take_fork(i);
        if (available((i+1))){
            take_fork((i + 1));
            eat();
        }else{
            put_fork(i);
        }
    }
}
```

Solution using Mutex

- Protect critical sections with a mutex
- Prevents deadlock
- But has performance issues
 - Only one philosopher can eat at a time

```
#define N 5
int forks = {1,2,3,4,5,1};

void philosopher(int i){
    while(TRUE){
        think(); // for some_time
        wait(mutex);
        take_fork(i);
        take_fork((i + 1));
        eat();
        put_fork(i);
        put_fork((i + 1));
        signal(mutex);
    }
}
```

Solution to Dining Philosophers

Uses N semaphores ($s[0]$, $s[1]$, ..., $s[N-1]$) all initialized to 0, and a mutex
Philosopher has 3 states: HUNGRY, EATING, THINKING

A philosopher can only move to EATING state if neither neighbor is eating

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

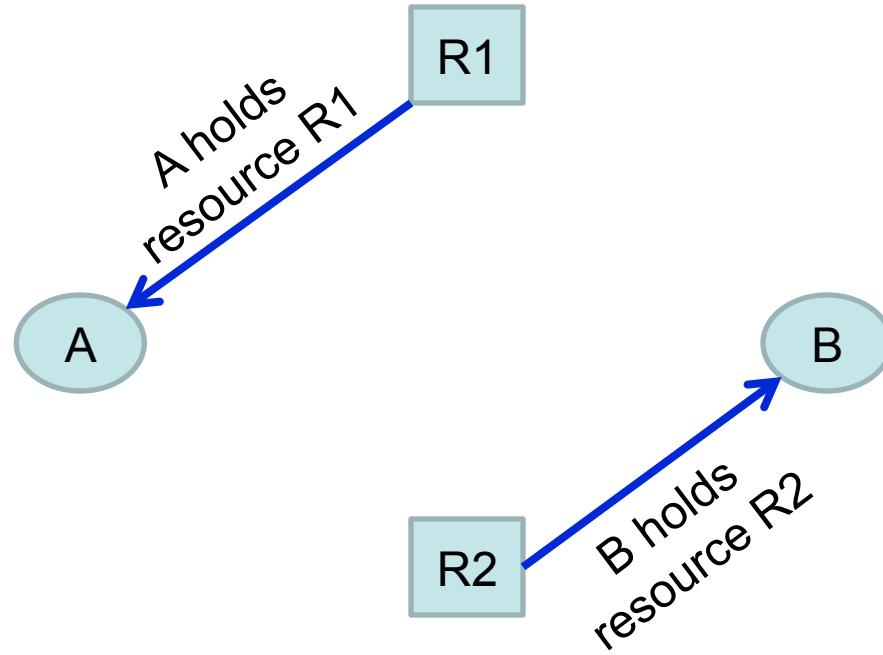
```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

Deadlocks

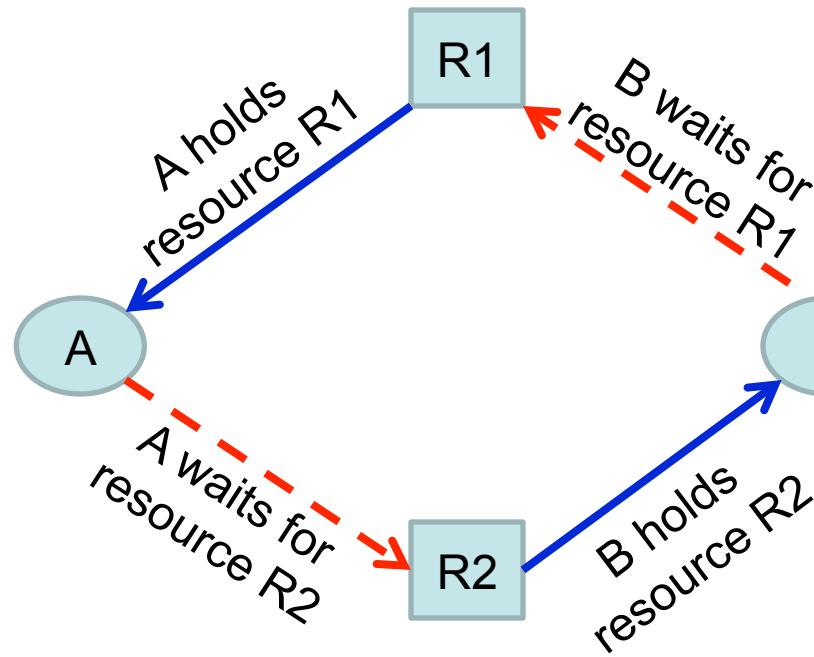
Deadlocks



Consider this situation:

deadlocks inherently have some kind of circular dependency that causes all the problems

Deadlocks



Resource Allocation Graph

RAGs are directed graphs and are used to model resource allocations
circles = processes
squares = resources
direction of arrow = direction of request
the requests are made to the OS which grants the requests to the processes and is responsible for managing them

A Deadlock Arises:

Deadlock : A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.

Conditions for Resource Deadlocks

1. Mutual Exclusion

- Each resource is either available or currently assigned to exactly one process

2. Hold and wait

- A process holding a resource, can request another resource

3. No preemption

- Resources previously granted cannot be forcibly taken away from a process only the process which holds a resource can release it, it cannot be snatched away from it

4. Circular wait

- There must be a circular chain of two or more processes, each of which is waiting for a resource held by the next member of the chain

IMPORTANT IDEA!!!!

All four of these conditions must be present for a resource deadlock to occur!!

deadlocks can be prevented by simply eliminating even one of these possible conditions completely

note that even when all the 4 are present
deadlocks are not compulsory
they are probabilistic and may not occur
at all or occur at low probabilities

Deadlocks :

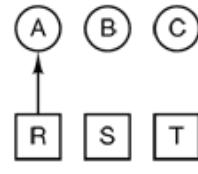
(A Chanced Event)

- Ordering of resource requests and allocations are probabilistic, thus deadlock occurrence is also probabilistic

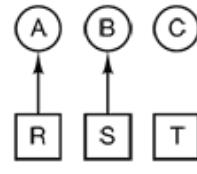
A	B	C
Request R		
Request S	Request S	
Release R	Request T	
Release S	Release S	Request T
(a)	(b)	(c)

1. A requests R
 2. B requests S
 3. C requests T
 4. A requests S
 5. B requests T
 6. C requests R
- deadlock

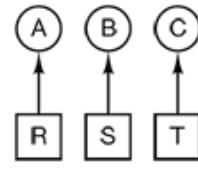
(d)



(e)

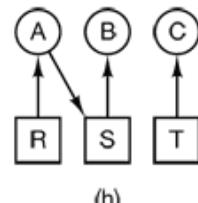


(f)

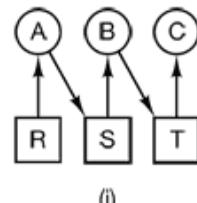


(g)

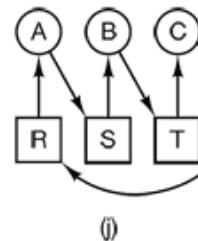
Deadlock occurs



(h)



(i)



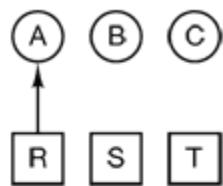
(j)

A	B	C
Request R	Request S	Request T
Request S	Request T	Request R
Release R	Release S	Release T
Release S	Release T	Release R
(a)	(b)	(c)

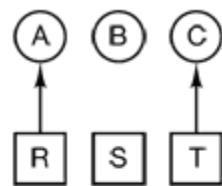
No dead lock occurrence
(B can be granted S
after step q)

1. A requests R
2. C requests T
3. A requests S
4. C requests R
5. A releases R
6. A releases S
no deadlock

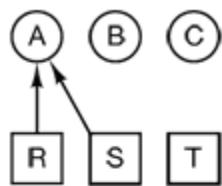
(k)



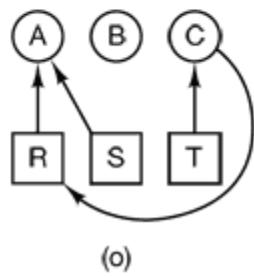
(l)



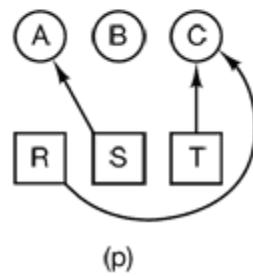
(m)



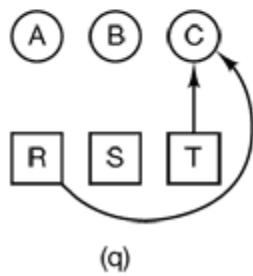
(n)



(o)



(p)



(q)

Should Deadlocks be handled?

- Preventing / detecting deadlocks could be tedious
- Can we live without detecting / preventing deadlocks?
 - What is the probability of occurrence?
 - What are the consequences of a deadlock? (How critical is a deadlock?)

deadlocks can be avoided by often having multiple resources - if two processes are requesting a printer, two printers can simply avoid the possibility of a deadlock

if the probability of a deadlock is very low and the consequences of a deadlock are not important, deadlocks can be ignored.
if there is a greater chance or a catastrophic outcome, then they must be detected/avoided or prevented

Handling Deadlocks

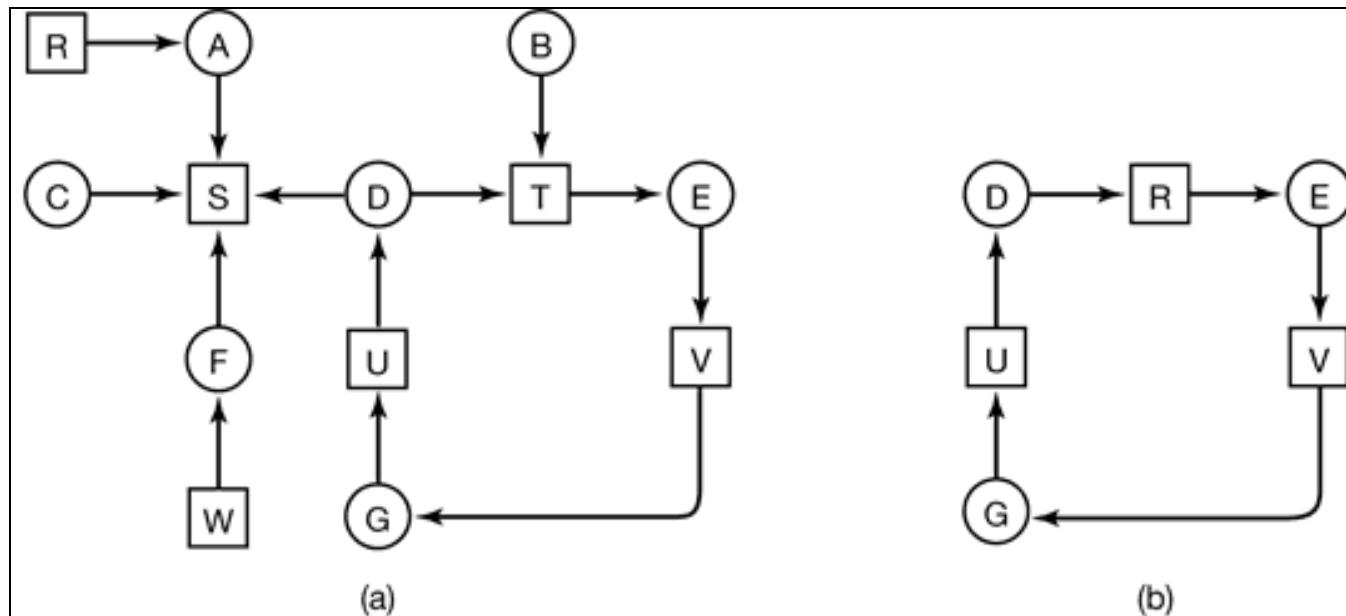
- Detection and Recovery
- Avoidance
- Prevention

Deadlock detection

- How can an OS detect when there is a deadlock?
- OS needs to keep track of
 - Current resource allocation
 - Which process has which resource
 - Current request allocation
 - Which process is waiting for which resource
- Use this information to detect deadlocks

Deadlock Detection

- Deadlock detection with **one resource of each type**
- Find cycles in resource graph



the idea of looking for cycles can only work when the system has one resource of each type

when none of the requests can be met, in any order, then there is a deadlocked scenario

Deadlock Detection

- Deadlock detection with multiple resources of each type

	Tape drives	Plotters	Scanners	CD Roms
E = (4	2	3	1)
Existing Resource Vector				

	Tape drives	Plotters	Scanners	CD Roms
A = (2	1	0	0)
Resources Available				

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

P₁ Current allocation matrix

$$P_2 \quad C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Current Allocation Matrix

Who has what!!

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

Request Matrix

Who is waiting for what!!

Process P_i holds C_i resources and requests R_i resources, where i = 1 to 3
Goal is to check if there is any sequence of allocations by which all current requests can be met. If so, there is no deadlock.

Deadlock Detection

- Deadlock detection with multiple resources of each type

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix}$$

Existing Resource Vector

	Tape drives	Plotters	Scanners	CD Roms
P ₁	4	2	3	1
P ₂	0	0	1	0
P ₃	2	0	0	1
P ₄	0	1	2	0

$$A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

Resources Available

	Tape drives	Plotters	Scanners	CD Roms
P ₁	2	1	0	0
P ₂	0	0	0	0
P ₃	0	0	0	0
P ₄	0	0	0	0

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

P₁ Current allocation matrix

P₂ Request matrix

P₃ P₁ cannot be satisfied

C = $\begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$

Current Allocation Matrix

P₂ P₂ cannot be satisfied

P₃ P₃ can be satisfied

R = $\begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$

Request Matrix

Process P_i holds C_i resources and requests R_i resources, where i = 1 to 3

Deadlock Detection

- Deadlock detection with multiple resources of each type

	Tape drives	Plotters	Scanners	CD Roms
$E = (4 \quad 2 \quad 3 \quad 1)$	4	2	3	1
Existing Resource Vector				

	Tape drives	Plotters	Scanners	CD Roms
$A = (2 \quad 1 \quad 0 \quad 0)$	2	1	0	0
Resources Available				

P_1 Current allocation matrix

$$P_2 \\ P_3 \\ C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Current Allocation Matrix

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

Request Matrix

P_3 runs and its allocation is $(2, 2, 2, 0)$

On completion it returns the available resources are $A = (4 \ 2 \ 2 \ 1)$

Either P_1 or P_2 can now run.

NO Deadlock!!!

Deadlock Detection

- Deadlock detection with multiple resources of each type

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix}$$

Existing Resource Vector

	Tape drives	Plotters	Scanners	CD Roms
E	4	2	3	1

$$A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

Resources Available

	Tape drives	Plotters	Scanners	CD Roms
A	2	1	0	0

$$\sum_{i=1}^n C_{ij} + A_j = E_j$$

$$P_1 \quad \text{Current allocation matrix}$$
$$P_2 \quad C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Current Allocation Matrix

$$P_3 \quad \text{Request matrix}$$
$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 1 & 0 \end{bmatrix}$$

Request Matrix

P_1 cannot be satisfied

P_2 cannot be satisfied

P_3 cannot be satisfied
deadlock

Process P_i holds C_i resources and requests R_i resources, where $i = 1$ to 3
Deadlock detected as none of the requests can be satisfied

Deadlock Recovery

What should the OS do when it detects a deadlock?

- Raise an alarm
 - Tell users and administrator
- Preemption
 - Take away a resource temporarily (frequently not possible)
- Rollback
 - Checkpoint states and then rollback
- Kill low priority process
 - Keep killing processes until deadlock is broken
 - (or reset the entire system)

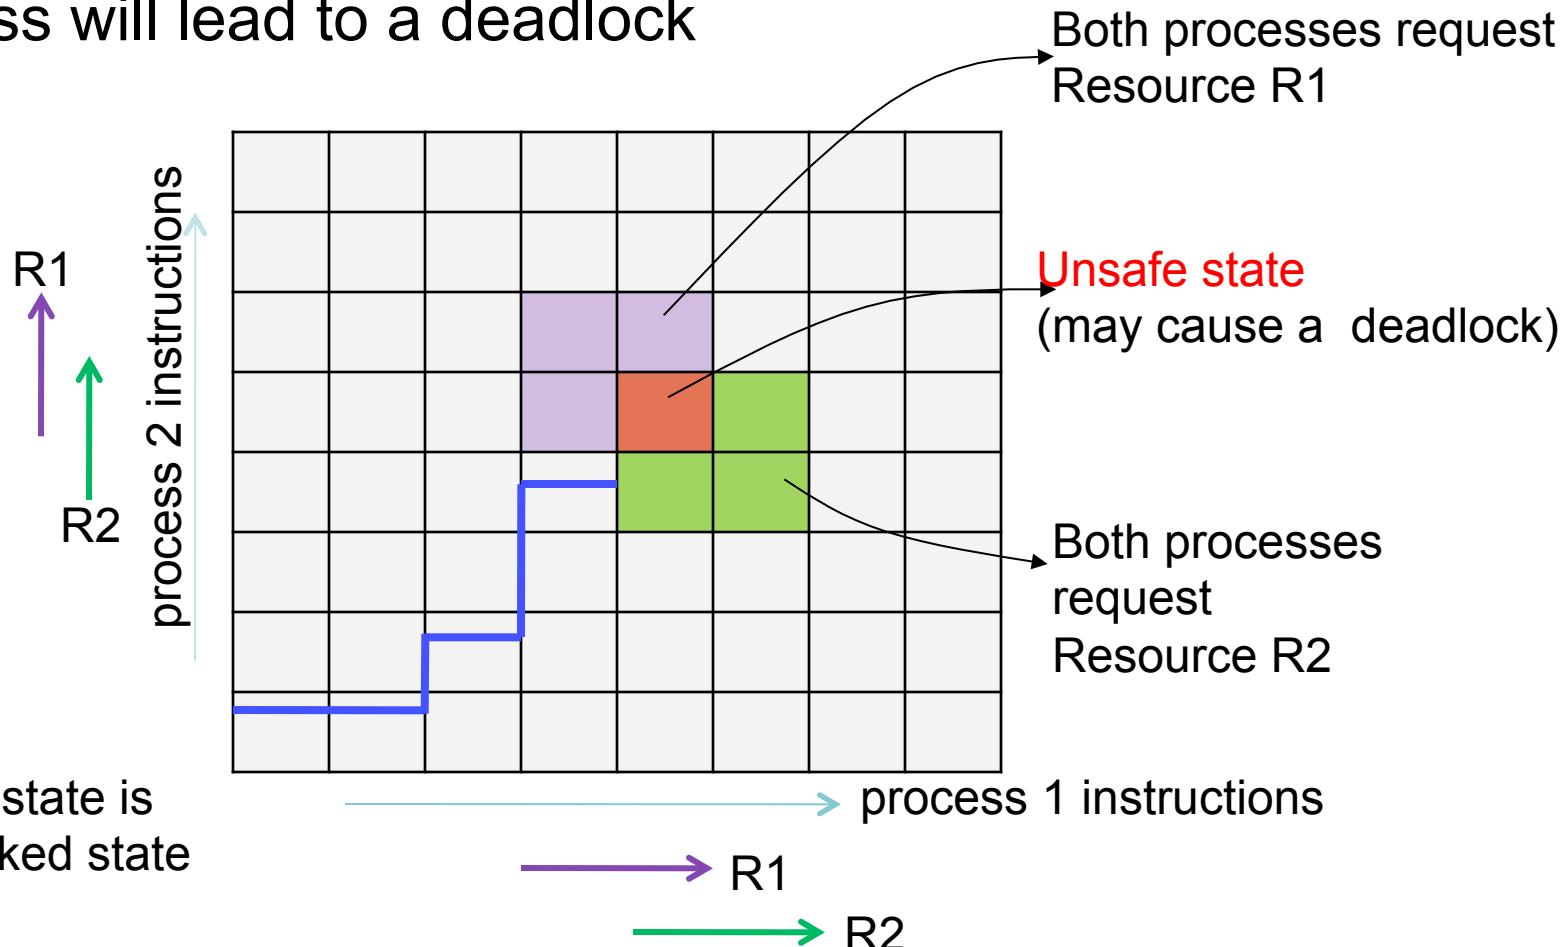
the checkpoint is when the complete PCB is stored onto the DISK to which it can return to when there is a deadlock

typically the process with a lower priority is killed until the deadlock dependency is broken

Deadlock Avoidance

evidently, the state where both the processes request access to both the resources is where there could potentially be a deadlock. deadlock avoidance scheduling ensures that a path is taken in the state space that avoids all possible unsafe states.

- System decides in advance if allocating a resource to a process will lead to a deadlock



Deadlock Avoidance

Is there an algorithm that can always avoid deadlocks by conservatively make the right choice.

- Ensures system never reaches an unsafe state
- **Safe state** : A state is said to be safe, if there is some scheduling order in which every process can run to completion even if all of them suddenly requests their maximum number of resources immediately
- An unsafe state **does not have to** lead to a deadlock; *it could lead to a deadlock*

Example with a Banker

- Consider a banker with 4 clients (P_1, P_2, P_3, P_4).
 - Each client has certain credit limits (totaling 20 units)
 - The banker knows that max credits will not be used at once, so he keeps only 10 units

	Has	Max
A	3	9
B	2	4
C	2	7

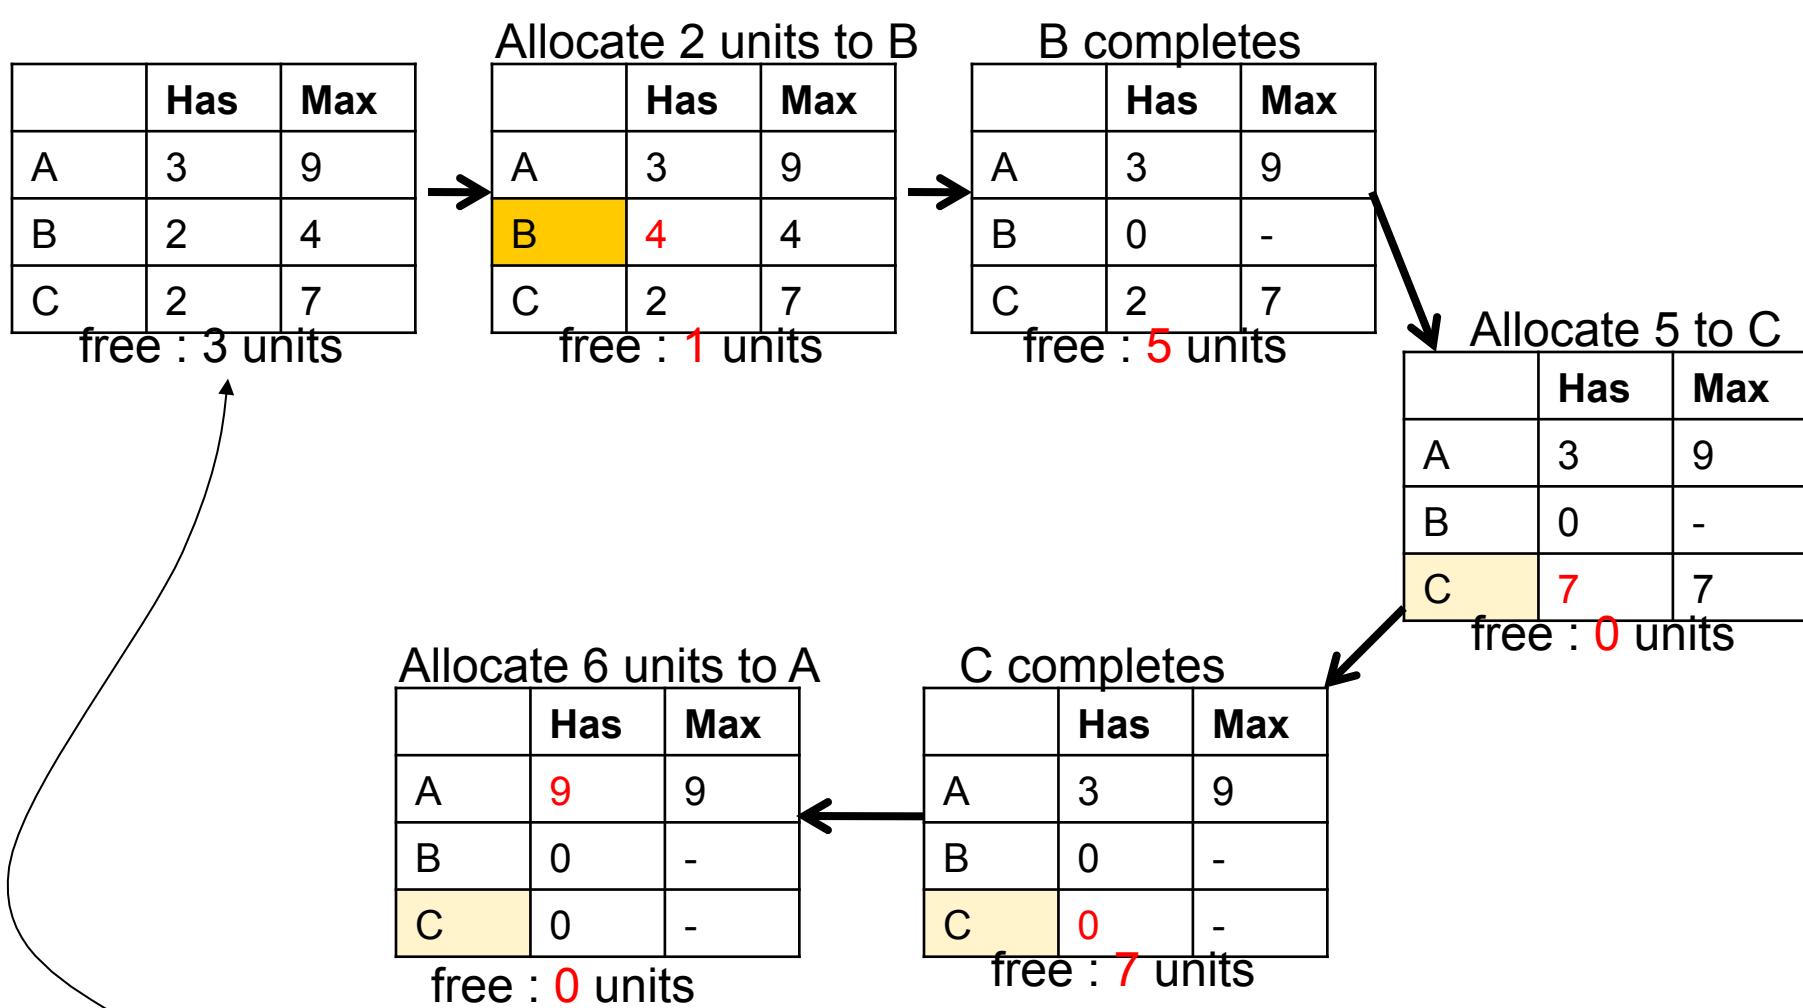
Total : 10 units

free : 3 units



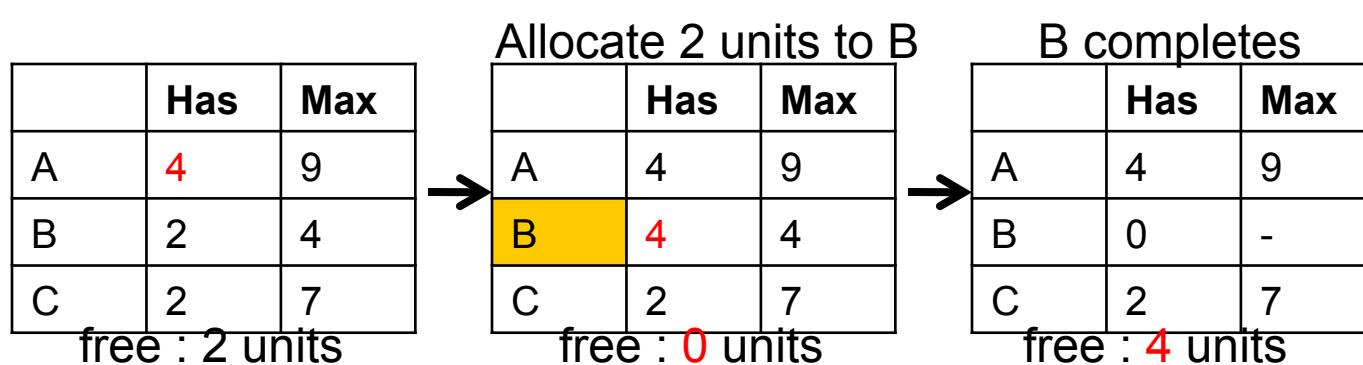
- Clients declare **maximum** credits in advance. The banker can allocate credits provided no unsafe state is reached.

Safe State



This is a safe state because there is some scheduling order in which every process executes

Unsafe State



This is an unsafe state because there exists NO scheduling order in which every process executes

Banker's Algorithm (with a single resource)

When a request occurs

- If(**is_system_in_a_safe_state**)
 - Grant request
- else
 - postpone until later



Please read Banker's Algorithm with multiple resources from Modern Operating Systems, Tanenbaum

Deadlock Prevention

- Deadlock avoidance not practical, need to know maximum requests of a process
- Deadlock prevention
 - Prevent at-least one of the 4 conditions
 1. Mutual Exclusion
 2. Hold and wait
 3. No preemption
 4. Circular wait

Prevention

1. Preventing Mutual Exclusion

- Not feasible in practice
- But OS can ensure that resources are optimally allocated

2. Hold and wait

- One way is to achieve this is to require all processes to request resources before starting execution
 - May not lead to optimal usage
 - May not be feasible to know resource requirements

3. No preemption

- Pre-empt the resources, such as by virtualization of resources (eg. Printer spools)

4. Circular wait

- One way, process holding a resource cannot hold a resource and request for another one
- Ordering requests in a sequential / hierarchical order.

Hierarchical Ordering of Resources

- Group resources into levels
(i.e. prioritize resources numerically)
- A process may only request resources at higher levels than any resource it currently holds
- Resource may be released in any order
- eg.
 - Semaphore s1, s2, s3 (with priorities in increasing order)
down(S1); down(S2); down(S3) ; → allowed
down(S1); down(S3); down(S2); →not allowed